

Amazon DynamoDB Deep Dive

Course Navigation

Course Introduction

Section 1

Databases 101

Section 2

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6



Linux Academy

Course Introduction

Getting Started

Course Navigation

Course Introduction

Section 1

Getting Started

Course Introduction

About the Training Architect

Course Features and Tools

Databases 101

Section 2

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6



Dear students,

Welcome to the Linux Academy **Amazon DynamoDB Deep Dive** course. This course will help you master DynamoDB!

In this course, you will learn the basics of DynamoDB and how it differs from traditional relational database management systems. A real-world scenario project will help guide you through each of the concepts presented.

This course is intended for all skill levels, even if you're totally new to database development.

By the end of this course, you'll have the necessary skills and knowledge to effectively build scalable, high performance applications using DynamoDB.

Thank you for taking the course. Let's get started!

Warm Regards,

Mark Richman

AWS Training Architect

Linux Academy

[Back to Main](#)



Linux Academy

Course Introduction

Section 1

Databases 101

Section 2

Database Management Systems

Relational DBs and SQL

NoSQL Database Models

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

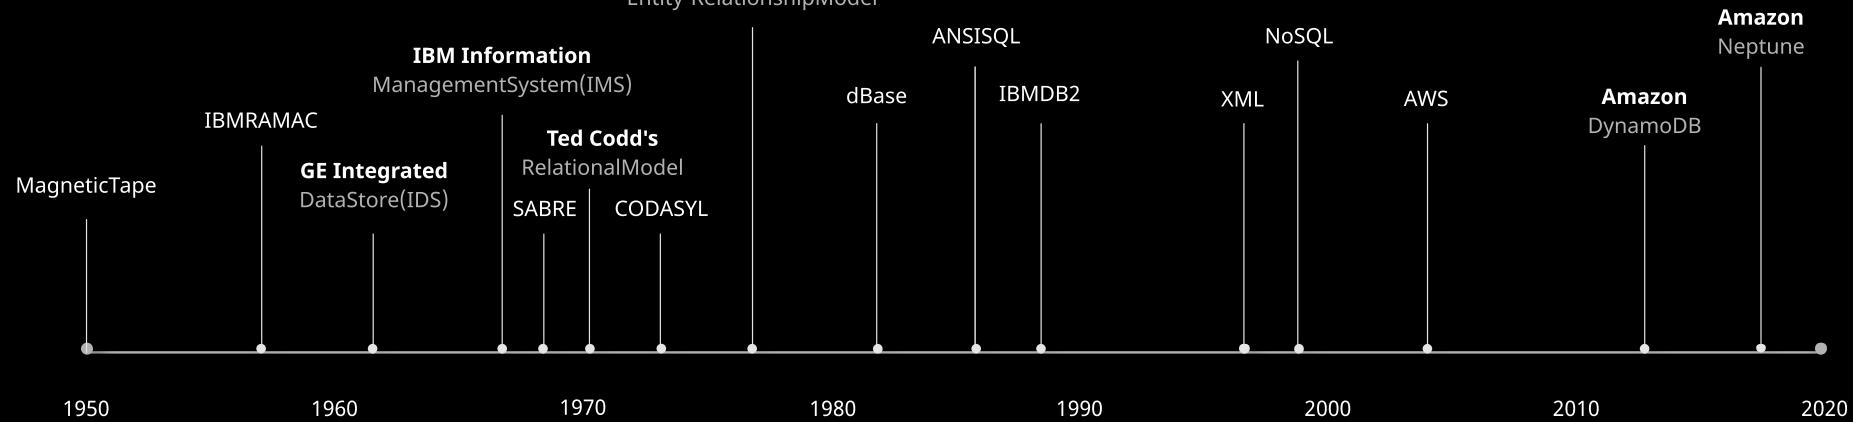
Section 5

Advanced DynamoDB

Section 6

History and Timeline

Peter Chen's
Entity-Relationship Model



Next

Back to Main

Course Introduction

Section 1

Databases 101

Section 2

Database Management Systems

Relational DBs and SQL

NoSQL Database Models

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

Microsoft Excel Is Not a Database ... or Is It?

Many people have typically used flat files or **spreadsheet software** as a type of database. In this example, orders are entered into a spreadsheet as they come in.

Over time, redundant data accumulates, along with **potentially contradicting values**.

For example, if Linda Oliver called in about her order, you would see three different entries on different dates, each with a different address.

This situation can easily lead to incorrect shipments, lost revenue, and unhappy customers.

How do we resolve this?

The screenshot shows a Microsoft Excel spreadsheet titled "Orders". The data is organized into columns: A (Name), B (Product), C (Quantity), D (Date), and E (Address). There are 11 rows of data, with row 1 serving as the header. The data includes multiple entries for Linda Oliver, demonstrating redundancy and potential contradictions. The Excel interface at the top includes tabs for Home, Insert, Draw, Page Layout, Formulas, etc., and a status bar at the bottom indicates "Ready".

	A	B	C	D	E
1	Name	Product	Quantity	Date	Address
2	Lakisha Holder	Colored Pencils	1	8/19	1626 Happy Hollow Road
3	Daniel Taylor	Ceramic Mug	2	8/19	3227 Mercer Street
4	Catherine Owens	Paper Towels	4	8/19	1857 Cedarstone Drive
5	Linda Oliver	AA Batteries	8	8/19	882 Oakwood Circle
6	Ralph Randolph	Earl Grey Tea	3	8/19	4182 Jim Rosa Lane
7	Terry Cobb	Laundy Detergent	5	8/19	945 Oliver Street
8	Linda Oliver	Diapers	7	8/28	123 Main Street
9	Rodney Lewis	Shampoo	9	8/19	4267 Meadow Drive
10	Linda Oliver	Toothpaste	1	8/28	8861 Lexington Avenue
11					

Spreadsheets are indeed databases in the same way flat text files are, but they lack some fundamental features, which we'll cover shortly.

[Back](#)[Next](#)[Back to Main](#)

Course Introduction

Section 1

Databases 101

Section 2

Database Management Systems

Relational DBs and SQL

NoSQL Database Models

DynamoDB and Scenario

Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

Relational Data

Instead of having just one massive spread sheet, you could separate the information into **different tables**. Each table represents an entity in your system.

Orders

ID	CustomerID	ProductID	Quantity	Date
76354	13579	238	1	\$4.99
76355	13580	239	2	\$13.98
76356	13581	240	1	\$13.95
76357	13582	241	3	\$35.97
76358	13583	242	2	\$22.77
76359	13584	243	1	\$34.95
76360	13585	244	2	\$13.90

Tables are linked, or related, to each other using various IDs for reference. This is what forms a **relational database**.

Customers

ID	Name	Address
13579	LakishaHolder	126HappyHollowRoad
13580	DanielTaylor	3227MercerStreet
13581	CatherineOwens	1857CedarstoneDrive
13582	LindaOliver	882OakwoodCircle
13583	RalphRandolph	4182JimRosaLane
13584	TerryCobb	945OliverStreet
13585	RodneyLewis	4267MeadowDrive

Products

ID	Name	Price
238	Colored Pencils	\$4.99
239	Ceramic Mug	\$6.99
240	Paper Towels	\$13.95
241	AA Batteries	\$11.99
242	Early Gray Tea	\$7.59
243	Laundry Detergent	\$14.89
244	Diapers	\$34.95
245	Shampoo	\$6.95

Back

Next

Course Introduction

Section 1

Databases 101

Section 2

Database Management Systems

Relational DBs and SQL

NoSQL Database Models

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

Terminology

Table: A collection of related data, organized in rows and columns

Row: A set of related values for all the columns in a given table. Often called a **record, item, or tuple.**

Column: A unit of named data that has a particular data type (e.g., number, text, or date)

Schema: Overall structure of the database, relationships between tables, and the rules that govern them

Key: Columns on which an index is constructed to allow rapid and/or sorted access to a table's row

- **Candidate Key:** Any set of columns that can uniquely identify a row
- **Primary Key:** A candidate key that may uniquely identify any given row
- **Foreign Key:** Uniquely identifies a value (primary key) in a different table

Query: A data retrieval statement that specifies the columns and tables from which data is to be retrieved; and optionally:

- Conditions the data must satisfy
- Computations to be performed on the retrieved column values
- A desired ordering of the result set
- **Examples:**
 - SELECT * FROM Customer WHERE ID = 76355
 - SELECT ID, Price FROM Product ORDER BY Price DESC
 - SELECT * FROM Order WHERE Date = 2019-04-02 AND Total > 100

[Back](#)

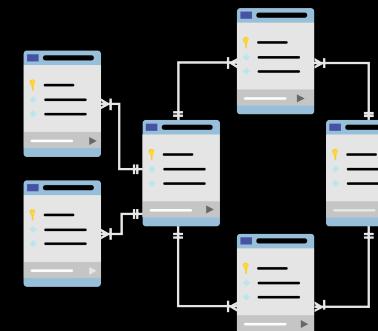
Columns				
ID	CustomerID	ProductID	Quantity	Total
76354	13579	238	1	\$4.99
76355	13580	239	2	\$13.98
76356	13581	240	1	\$13.95
76357	13582	241	3	\$35.97
76358	13583	242	2	\$22.77
76359	13584	243	1	\$34.95
76360	13585	244	2	\$13.90

Rows

Primary Key

Foreign Key

Foreign Key



Schema

[Back to Main](#)[Next](#)

Course Introduction

Section 1

Databases 101

Section 2

Database Management Systems

Relational DBs and SQL

NoSQL Database Models

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

Database Components

1

Hardware

Physical electronic devices such as computers, I/O devices, storage devices, etc. This provides the **interface** between computers and real-world systems.

2

Software

Programs used to **control and manage the overall database**. This includes the DBMS software itself, operating system, network software used to share the data among users, and the application programs used to access data in the DBMS.

3

Data

The most important component of the **DBMS**. Data can be either the actual operational data or metadata.

4

Data Access Method

Programming language or similar commands to access, query, and **manipulate the data** in the DBMS.

5

Users

People who need the information from the database to carry out their primary **business responsibilities**.

Back

Back to Main

Course Introduction

Section 1

Databases 101

Section 2

Database Management Systems

Relational DBs and SQL

NoSQL Database Models

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

ACID Properties in Relational Databases

A sequence of database operations that satisfies these properties is called a **transaction**.

In order to guarantee validity of data after a transaction, operations must satisfy these properties:

A

Atomicity

A sequence of **operations** where either all occur or nothing occurs.

C

Consistency

The state of the **database**, both before and after execution of an operation, remains consistent (i.e., free of any data integrity errors) whether or not the transaction commits or is rolled back.

I

Isolation

Changes made by a transaction are **isolated from the rest** of the system until after the transaction has committed.

D

Durability

Committed data is saved by the system such that, even in the event of a failure and system restart, the data is available in its correct state.

Next

Back to Main

Course Introduction

Section 1

Databases 101

Section 2

Database Management Systems**Relational DBs and SQL**

NoSQL Database Models

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

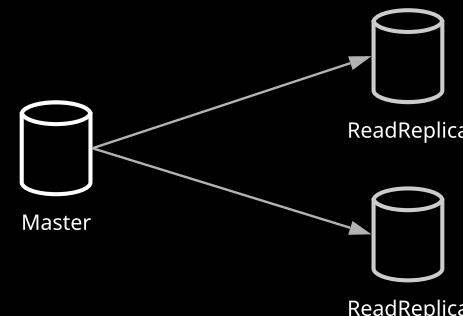
Advanced DynamoDB

Section 6

SQL Databases and Scaling

Replication

Too much load



Sharding

Too much data

**Pros:**

- Offloads master node
- Read replicas can scale horizontally

Cons:

- Does not increase the speed of writes
- Replication lag

Pros:

- Can store larger data sets
- Can handle more load than a single master node

Cons:

- Queries become more complex
 - Ex: Range queries might hit all nodes
- Joins become difficult — your tables are no longer all on the same node

[Back](#)[Next](#)[Back to Main](#)

Course Introduction

Section 1

Databases 101

Section 2

Database Management Systems

Relational DBs and SQL

NoSQL Database Models

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

[Back to Main](#)**Partitioning**

User



User



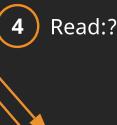
2 Network disconnect or packet loss



Database Node A



Database Node B



Course Introduction

Section 1

Databases 101

Section 2

Database Management Systems
Relational DBs and SQL

NoSQL Database Models

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

CAP Theorem

How do we overcome the limitations of **SQL databases**? Scaling horizontally comes with tradeoffs ...

The **CAP theorem** states that it is impossible for a distributed data store to simultaneously provide more than two out of the following three guarantees:

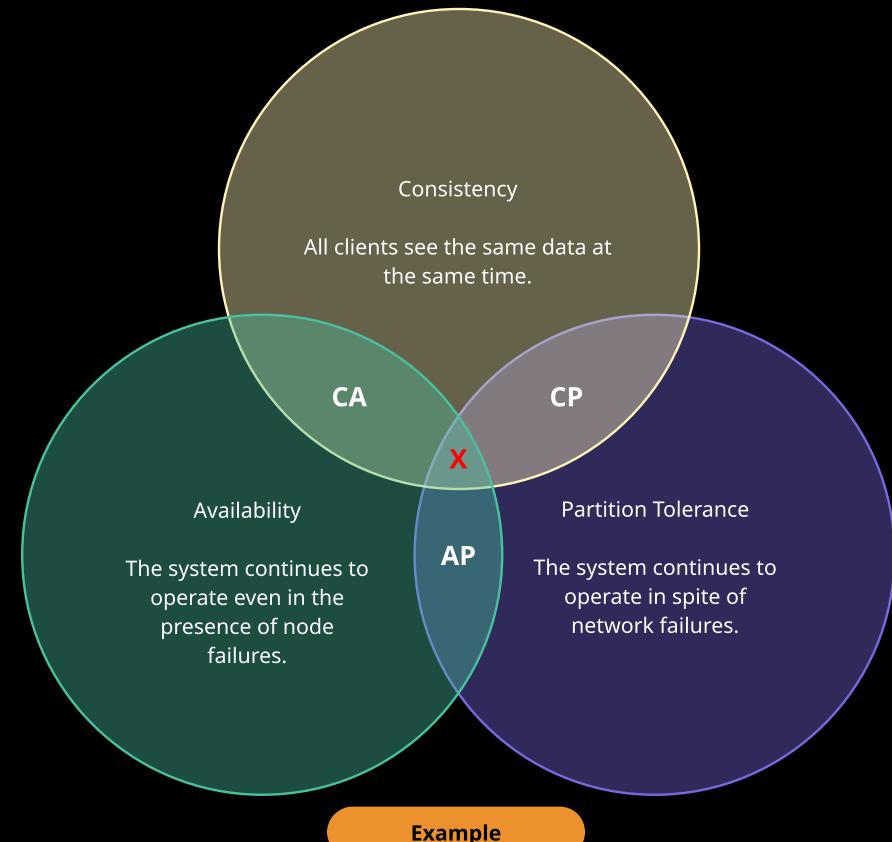
- **CA**: Data is consistent between all nodes — as long as all nodes are online — and you can read/write from any node and be sure the data is the same, but if you ever develop a partition between nodes, the data will be out of sync (and won't re-sync once the partition is resolved).
- **CP**: Data is consistent between all nodes and maintains partition tolerance (preventing data de-sync) by becoming unavailable when a node goes down.
- **AP**: Nodes remain online even if they can't communicate with each other and will re-sync data once the partition is resolved, but you aren't guaranteed that all nodes will have the same data (either during or after the partition).

In a **distributed system**, we always need to select partition tolerance; we need to tolerate packet loss on the network level. So in reality, we only have a choice between consistency and availability (CA).

Given disconnected partitions (packet loss), a database node can either **return an error or stale data**.

SQL databases select both C and A. They suffer from partition intolerance and generally run on only one machine. Replication technologies do not select both C and A.

NoSQL databases select either C or A.

[Back](#)[Back to Main](#)

Course Introduction

Section 1

Databases 101

Section 2

Database Management Systems

Relational DBs and SQL

NoSQL Database Models

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

Introduction to NoSQL

NoSQL databases are built for specific data models and offer flexible schemas. These databases are well suited to modern applications such as mobile, web, and gaming that require high scalability, flexibility, and performance.

SQL

NoSQL

Optimized for storage	Optimized for compute
Normalized/relational	Denormalized/hierarchical
Ad-hocqueries	Instantiatedviews
Scalevertically	Scalehorizontally
GoodforOLAP	BuiltforOLTPatscale

Next

Back to Main

Course Introduction

Section 1

Databases 101

Section 2

Database Management Systems

Relational DBs and SQL

NoSQL Database Models

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

NoSQL Engine Types

Key-value: Uses a simple key-value method to store data. Stored as a collection of key-value pairs, where the key serves as a unique identifier. Schema is defined per item.

Use cases: Session store, shopping cart

Document-oriented: Designed to store JSON-like data

Use cases: Content management, catalogs

Column-oriented: Stores each column of data in sequential blocks on disk or in memory. Ideal for analytical queries that perform aggregate operations (sum, avg, etc.) over a small number of columns. **Takes advantage of compression**, as similar data is stored together. Slow for single record operations.

Use cases: Analytics, data warehouses

Graph: Designed to store relationships. Data is stored as nodes and edges. Very fast because relationships are not calculated at query time but persisted in the database.

Use cases: Fraud detection, recommendation engines

Back

ProductID	Attributes
1234	Book,Odyssey,Homer
3284	Album,6Partitas,Bach
5283	Movie,Dune,Lynch
9833	Movie,StarWars,Lucas

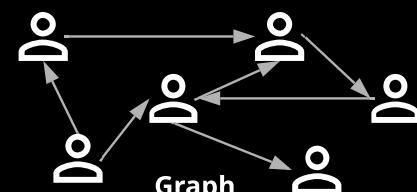
Key-Value

Date	Price	Size
2019-01-20	10.1	10
2019-01-21	10.3	20
2019-01-22	10.5	40
2019-01-23	11.2	3

Column-Oriented

```
[ {  
  year: 2020,  
  title: 'Dune',  
  info: {  
    directors: ['Denis Villeneuve', 'Not David Lynch'],  
    release_date: '2020-12-18T00:00:00Z',  
    rating: 9.2,  
    genres: ['Fantasy', 'Sci-Fi'],  
    image_url: 'http://example.com/images/foo.jpg',  
    plot: 'Paul Atreides rides worms around Arrakis.',  
    actors: ['Timothée Chalamet', 'Zendaya']  
  }  
}
```

Document-Oriented



Next

Course Introduction

Section 1

Databases 101

Section 2

Database Management Systems

Relational DBs and SQL

NoSQL Database Models

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

ACID vs. BASE

Basic Availability: The database appears to work most of the time.

Soft state: Stores don't have to be write-consistent, nor do different replicas have to be mutually consistent all the time.

Eventual consistency: Stores exhibit consistency at some later point (e.g., lazily at read time).

The main difference between **ACID** and **BASE** compliance is that ACID ensures that at the point in time of the transaction, ACID compliance is respected. BASE compliance allows for that compliance to be violated for a time as long as it eventually gets to a compliant end state.

We can think of ACID and BASE as a **continuum** instead of mutually exclusive choices:

ACID

- Strong consistency
- Isolation
- Focus on "commit"
- Nested transactions
- Availability
- Conservative/pessimistic
- Difficult to evolve schema

BASE

- Weak consistency (stale data acceptable)
- Availability first
- Best effort
- Approximate answers are acceptable
- Aggressive/optimistic
- Faster and easier schema evolution



Continuum of choices ...

Back

Next

Back to Main

Course Introduction

Section 1

Databases 101

Section 2

Database Management Systems

Relational DBs and SQL

NoSQL Database Models

DynamoDB and Scenario Introduction

Section 3

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

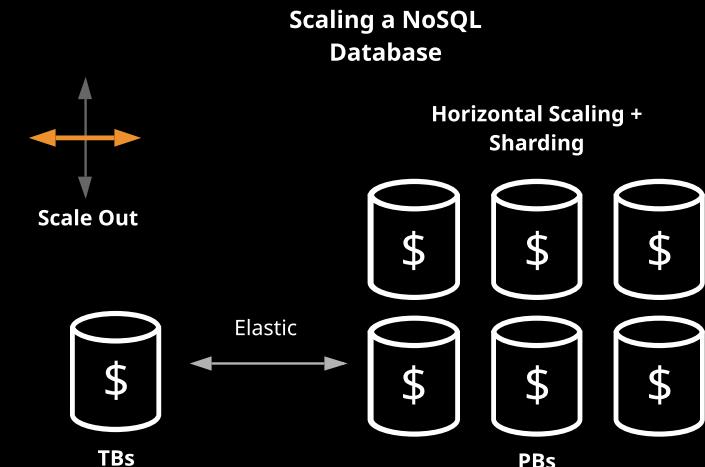
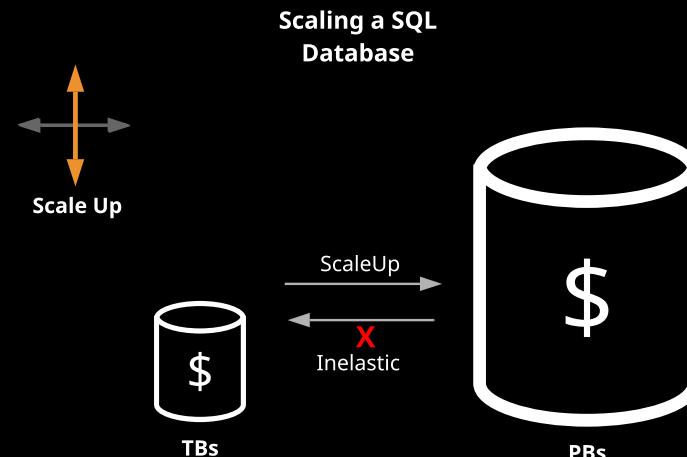
Section 5

Advanced DynamoDB

Section 6

How Scaling in NoSQL Is Different

The biggest advantage of NoSQL databases is horizontal scaling or **sharding**. In a scale-out architecture, a distributed set of nodes known as a **cluster** is used as the basic architecture. It provides highly elastic scaling capability, enabling you to add nodes to handle load on the fly. This is the opposite of a scale-up architecture associated with **SQL databases**, which adds more resources to a single, larger machine.



Back

Back to Main

DynamoDB and Scenario Introduction

Course Navigation

DynamoDB 10,000 FT

Course Introduction

Section 1

Databases 101

Section 2

DynamoDB and Scenario Introduction

Section 3

DynamoDB 10,000 FT

Introducing Pinehead Records

Scenario Discussion and Architecture Thoughts

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

A Little History

2004

Oracle configured with clustering and replication fails under load during the holiday season.

Analysis reveals 70% of **DB operations** were key-value and 20% operated on a single table.

Amazon begins development on Dynamo, a purpose-built, highly available key-value data store.

2007

The Amazon engineers behind **Dynamo** publish Dynamo: Amazon's Highly Available Key-value Store (a.k.a., the "Dynamo Paper"), summarizing their learnings.

Dynamo is poorly received inside Amazon; engineers continue to use **S3 and SimpleDB services**.

Amazon redesigns Dynamo for ease of management, performance, high availability, security, and scalability.

2012

Amazon DynamoDB **launches**.

SimpleDB is **effectively deprecated**.

Next

Back to Main



Linux Academy

DynamoDB and Scenario Introduction

Course Navigation

DynamoDB 10,000 FT

Course Introduction

Section 1

Databases 101

Section 2

DynamoDB and Scenario Introduction

Section 3

DynamoDB 10,000 FT

Introducing Pinehead Records

Scenario Discussion and Architecture Thoughts

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

What is DynamoDB?

Here's the marketing pitch...



Amazon DynamoDB

- NoSQL managed database service
- Supports both key-value and document data models
- It's really fast
 - Consistent responsiveness
 - Single-digit millisecond
- Unlimited throughput and storage
- Automatic scaling up or down
- Handles **trillions** of requests per day
- ACID transaction support
- On-demand backups and point-in-time recovery
- Encryption at rest
- Data is replicated across multiple Availability Zones
- Service-level agreement (SLA) up to **99.999%**

Back

Next

Back to Main



Linux Academy

DynamoDB and Scenario Introduction

Course Navigation

DynamoDB 10,000 FT

Course Introduction

Section 1

Databases 101

Section 2

DynamoDB and Scenario Introduction

Section 3

DynamoDB 10,000 FT

Introducing Pinehead Records

Scenario Discussion and Architecture Thoughts

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

Where Does DynamoDB Fit In?



Amazon Relational Database Service (RDS)

Support for Amazon Aurora, PostgreSQL, MySQL, MariaDB, Oracle Database, and SQL Server



Amazon DynamoDB

Key-value and document database



Amazon ElastiCache

Managed, Redis- or Memcached-compatible in-memory data store



Amazon Neptune

Graph database for applications that work with highly connected data sets



Amazon Redshift

Petabyte-scale data warehouse service



Amazon QLDB

Ledger database providing a cryptographically verifiable transaction log



Amazon DocumentDB

MongoDB-compatible database service

Back

Back to Main



Linux Academy

DynamoDB and Scenario Introduction

Course Navigation

Introducing Pinehead Records

Course Introduction

Section 1

Databases 101

Section 2

DynamoDB and Scenario Introduction

Section 3

DynamoDB 10,000 FT

Introducing Pinehead Records

Scenario Discussion and Architecture Thoughts

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

Example Application Scenario

Pinehead Records is the world's largest retailer of vinyl records. Since establishing an online presence, high traffic and scaling issues have crippled the website on several occasions.

The current platform is Python, Flask, and MySQL.

The website must be migrated off its legacy relational database back end and onto a stable, high-performance platform that can scale with minimal operational overhead.

Application Migration Plan

Version 0: Legacy Relational Model

- Relational model in MySQL
- Limited database optimizations
- Limited application-level caching
- No indexes
- Inefficient queries
- Images (album cover art) stored on local filesystem
- User accounts in database

Version 1: Fundamental DynamoDB

- Naive migration from MySQL to DynamoDB
- Three DynamoDB tables mimicking the relational structure
- Images are moved to S3 with URI in a DynamoDB attribute
- No indexes
- User accounts in DynamoDB

Next

Back to Main



Linux Academy

DynamoDB and Scenario Introduction

Introducing Pinehead Records

Course Navigation

Course Introduction

Section 1

Databases 101

Section 2

DynamoDB and Scenario Introduction

Section 3

DynamoDB 10,000 FT

Introducing Pinehead Records

Scenario Discussion and Architecture Thoughts

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

Version 2: Intermediate DynamoDB

- Some optimizations
- Better table structure (single hierarchical table)
- Indexes
- Transactions
- User accounts in DynamoDB

Version 3: Advanced DynamoDB

- Federated web identity (Cognito)
- Fine-grained policies
- Lambda triggers
- Improved security
- DynamoDB Accelerator (DAX)

Back

[Back to Main](#)



Linux Academy

DynamoDB and Scenario Introduction

Course Navigation

Course Introduction

Section 1

Databases 101

Section 2

DynamoDB and Scenario Introduction

Section 3

DynamoDB 10,000 FT

Introducing Pinehead Records

Scenario Discussion and Architecture Thoughts

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

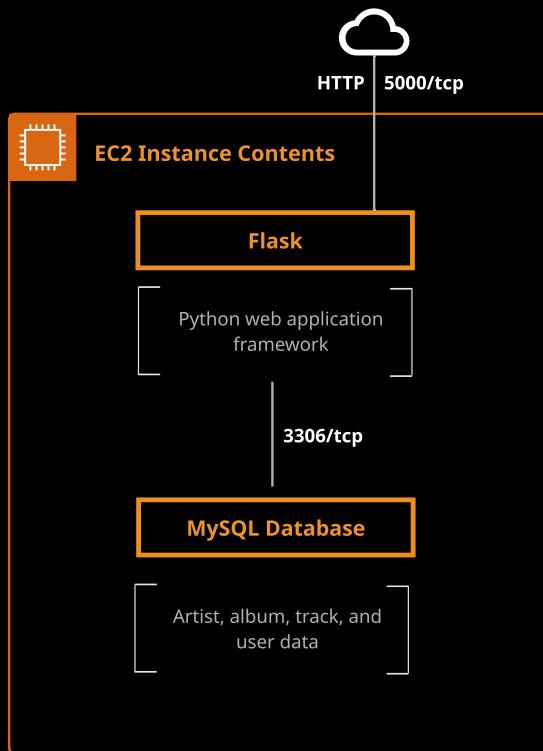
Advanced DynamoDB

Section 6

Scenario Discussion and Architecture Thoughts

Demo Application Architecture

Version 0



Next

Back to Main



Linux Academy

DynamoDB and Scenario Introduction

Course Navigation

Course Introduction

Section 1

Databases 101

Section 2

DynamoDB and Scenario Introduction

Section 3

DynamoDB 10,000 FT

Introducing Pinehead Records

Scenario Discussion and Architecture Thoughts

DynamoDB Fundamentals

Section 4

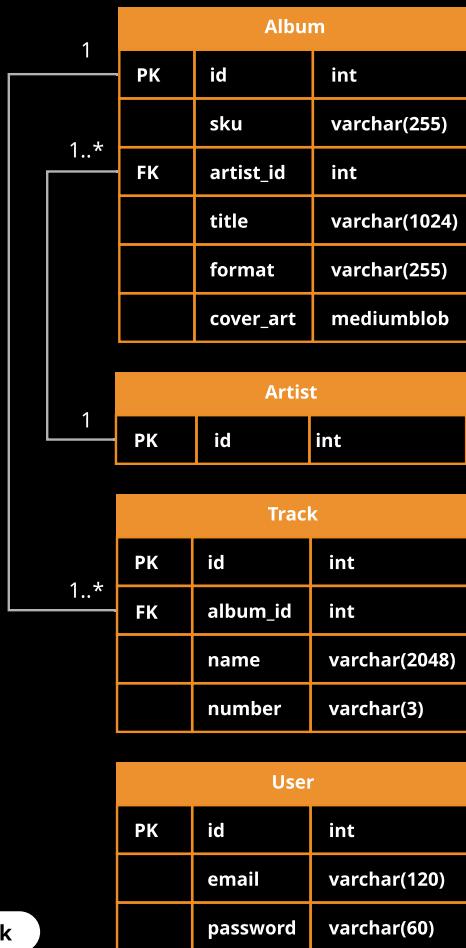
Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

Scenario Discussion and Architecture Thoughts



Back

Next

Back to Main



Linux Academy

DynamoDB and Scenario Introduction

Scenario Discussion and Architecture Thoughts

Course Navigation

Course Introduction

Section 1

Databases 101

Section 2

DynamoDB and Scenario Introduction

Section 3

DynamoDB 10,000 FT

Introducing Pinehead Records

Scenario Discussion and Architecture Thoughts

DynamoDB Fundamentals

Section 4

Intermediate DynamoDB

Section 5

Advanced DynamoDB

Section 6

[Back to Main](#)

Why DynamoDB?

1

Learning something new is hard. I already know SQL.

2

DynamoDB has no query language.

3

I have to choose between SQL and DynamoDB.

4

We already have a site license for *<commercial database>*.

5

It's just as hard to maintain my code with DynamoDB as it is with SQL, so what's the point?

6

My database isn't that big, isn't DynamoDB for big databases?

7

I enjoy writing SQL.

8

Why do I suddenly need a new database technology? Is this a gimmick?

[Back](#)



Linux Academy

DynamoDB Fundamentals

Tables and Item Architecture

Course Navigation

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Select the items below regarding tables and item architecture:

DynamoDB Architecture

DynamoDB Items

Previous Section

Next

Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

DynamoDB Architecture

Course Navigation

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

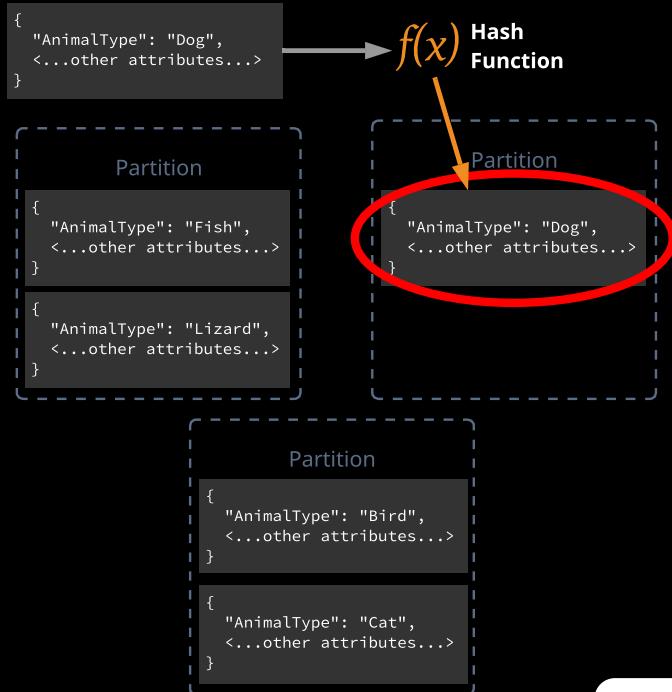
Data Model - Version 1

Version 1 Data Model Overview

Partitions and Data Distribution

DynamoDB stores data in partitions. A partition is an allocation of storage for a table, backed by solid-state drives (SSDs) and automatically replicated across multiple Availability Zones within an AWS Region.

To get the most out of DynamoDB throughput, create tables where the partition key has a large number of distinct values. Applications should request values fairly uniformly and as randomly as possible.



Next

Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

DynamoDB Architecture

Course Navigation

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Table: Collection of data. DynamoDB tables must contain a name, primary key, and the required read and write throughput values. **Unlimited** size.

Partition Key: A simple primary key, composed of one attribute known as the **partition key**. This is also called the **hash attribute**.

Partition and Sort Key: Also known as a **composite primary key**, this type of key comprises two attributes. The first attribute is the **partition key**, and the second attribute is the **sort key**. This is also called the **range attribute**.

- Mandatory
- Key-value access pattern
- Determines data distribution



- Optional
- Model 1-to-many relationships
- Enables rich query capabilities
 - All items for key
 - ==, <, >, >=, <=
 - Begins with, between, etc.
 - Sorted results and counts

Back

Next

Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

DynamoDB Performance

On-Demand Capacity

- Database scales according to demand
- Good option for:
 - New tables with unknown workloads
 - Applications with unpredictable traffic
 - Prefer to pay as you go

Provisioned Capacity

- Allows us to have **consistent and predictable** performance
- Specify expected read and write throughput requirements
- Read Capacity Units (**RCU**)
- Write Capacity Units (**WCU**)
- Price is determined by provisioned capacity
- Cheaper per request than On-Demand mode
- **Good option for:**
 - Applications with predictable traffic
 - Applications whose traffic is consistent or ramps gradually
 - Capacity requirements can be forecasted, helping to control costs

Both capacity modes have a limit of **40,000 RCUs and 40,000 WCUs**.

You can switch between modes only once per **24 hours**.

Table Addressing via ARN

arn:aws:dynamodb:us-east-1:123456789012:table/album

Region

Account

Table

Back



DynamoDB Fundamentals

DynamoDB Items

Course Navigation

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

- DynamoDB Architecture
- DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

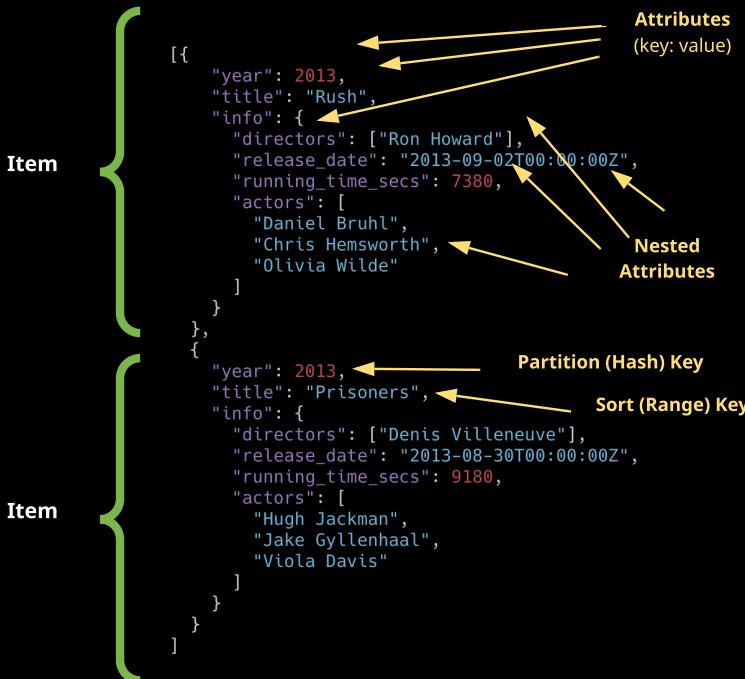
Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Item: A table may contain multiple items. An item is a unique group of attributes. Items are similar to rows or records in a traditional relational database. Items are limited to 400 KB.

Attribute: Fundamental data element. Similar to fields or columns in an RDBMS.



Next

Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

DynamoDB Items

Course Navigation

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Data Types

Scalar: Exactly one value — number, string, binary, boolean, and null.

Applications must encode binary values in base64-encoded format before sending them to DynamoDB.

Document: Complex structure with nested attributes (e.g.. JSON) — list and map.

Document Types

List: Ordered collection of values

FavoriteThings: ["Cookies", "Coffee", 3.14159]

Map: Unordered collection of name-value pairs (similar to JSON)

```
{  
    Day: "Monday",  
    UnreadEmails: 42,  
    ItemsOnMyDesk: [  
        "Coffee Cup",  
        "Telephone",  
        {  
            Pens: { Quantity : 3},  
            Pencils: { Quantity : 2},  
            Erasers: { Quantity : 1}  
        }  
    ]  
}
```

Set: Multiple scalar values of the same type — string set, number set, binary set.

["Black", "Green", "Red"]

[42.2, -19, 7.5, 3.14]

["U3Vubnk=", "UmFpbnk=", "U25vd3k="]

Back

Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

Accessing DynamoDB

Course Navigation

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Select the items below regarding DynamoDB accessibility:

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Previous Section

Next

Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

Using the AWS Management Console

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Creating a Table

- **Table names** must be unique per AWS account and region.
 - Between 3 and 255 characters long
 - UTF-8 encoded
 - Case-sensitive
 - Contain a-z, A-Z, 0-9, _ (underscore), - (dash), and . (dot)
- **Primary key** must consist of a partition key or a partition key and sort key.
 - Only string, binary, and number data types are allowed for partition or sort keys
- **Provisioned capacity mode** is the default (free tier).
 - For provisioned capacity mode, read/write throughput settings are required
- **Secondary indexes** creates a local secondary index.
 - Must be created at the time of table creation
 - Same partition key as the table, but a different sort key
- **Provisioned capacity** is set at the table level.
 - Adjust at any time or enable auto scaling to modify them automatically
 - On-demand mode has a default upper limit of 40,000 RCU/WCU — unlike auto scaling, which can be capped manually

Next

Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

Using the AWS Management Console

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

DynamoDB Console Menu Items

- Dashboard
- Tables
 - Storage size and item count are **not** real-time (updated ~6h)
 - **Items:** Manage items and perform queries and scans.
 - **Metrics:** Monitor CloudWatch metrics.
 - **Alarms:** Manage CloudWatch alarms.
 - **Capacity:** Modify a table's provisioned capacity.
 - Free tier allows 25 RCU, 25 WCU, and 25 GB for 12 months
 - Linux Academy **Cloud Sandbox** within the Cloud Playground
 - **Indexes:** Manage global secondary indexes.
 - **Global Tables:** Multi-region, multi-master replicas
 - **Backups:** On-demand backups and point-in-time recovery
 - **Triggers:** Manage triggers to connect DynamoDB streams to Lambda functions.
 - **Access control:** Set up fine-grained access control with web identity federation.
 - **Tags:** Apply tags to your resources to help organize and identify them.
- Backups
- Reserved capacity
- Preferences
- DynamoDB Accelerator (DAX)

Back

Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

Using the AWS CLI

Course Navigation

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Installing the AWS CLI

- Preinstalled on Amazon Linux and Amazon Linux 2
- Linux Academy **Cloud Sandbox** within the Cloud Playground

Obtaining IAM Credentials

- Option 1: Create IAM access keys in your own AWS account.
- Option 2: Use Cloud Sandbox credentials.
- Note the access key ID and secret access key.

Configuring the AWS CLI

- aws configure
- aws sts get-caller-identity
- aws dynamodb help

Using DynamoDB with the AWS CLI

- aws dynamodb create-table
- aws dynamodb describe-table
- aws dynamodb put-item
- aws dynamodb scan

[Back to Main](#)

[Section 5](#)



Linux Academy

DynamoDB Fundamentals

Course Navigation

SDK Options

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

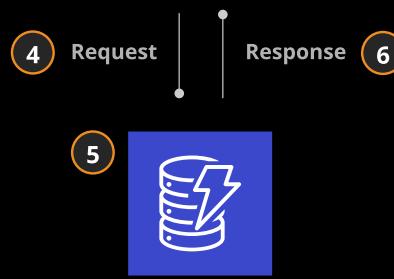
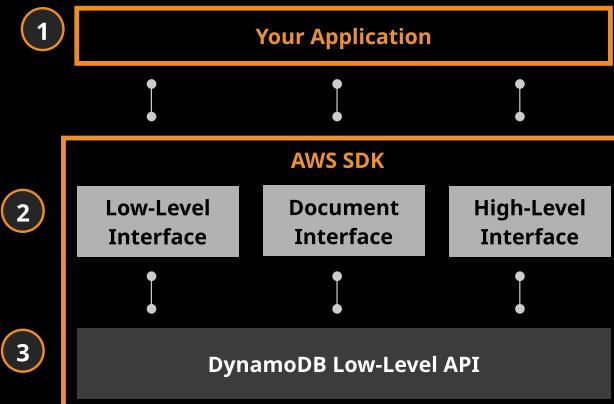
Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview



Back to Main

Section 5



Linux Academy

Next

DynamoDB Fundamentals

Course Navigation

SDK Options

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Data Type Descriptors



S – String

M – Map

N – Number

L – List

B – Binary

SS – String Set

BOOL – Boolean

NS – Number Set

NULL – Null

BS – Binary Set

Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

Course Navigation

SDK Options

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Low-Level API

Request

```
POST / HTTP/1.1
Host: dynamodb.us-east-1.amazonaws.com
Accept-Encoding: identity
Content-Length: 158
User-Agent: aws-cli/1.16.225
Content-Type: application/x-amz-json-1.0
Authorization: AWS4-HMAC-SHA256 Credential=...
X-Amz-Date: 20190916T141902Z
X-Amz-Target: DynamoDB_20120810.GetItem
```

```
{
  "TableName": "Music",
  "Key": {
    "Artist": {"S": "Queen"},
    "Year": {"N": "1989"}
  }
}
```

Data Type Descriptors

Response

```
HTTP/1.1 200 OK
x-amzn-RequestId: 6B5D43RPJ1PH5CS15RJUV66DFJVV4KQNS05AEMVJF66Q
x-amz-crc32: 3651145634
Content-Type: application/x-amz-json-1.0
Content-Length: 188
Date: 20190916T141903Z
```

```
{
  "Item": {
    "Artist": {"S": "Queen"},
    "Year": {"N": "1989"},
    "AlbumTitle": {"S": "The Miracle"},
    "TrackCount": {"N": "13"},
    "Barcode": {"S", "0077779235728"}
  }
}
```

Back

Next

Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

SDK Options

Course Navigation

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

Putitem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Document Interface

- Perform data plane operations (create, read, update, delete) on tables and indexes
- No need for data type descriptors
- Available in the AWS SDKs for Java, .NET, Node.js, and JavaScript in the browser

```
var docClient = new AWS.DynamoDB.DocumentClient()
var table = 'Music'
var artist = 'Queen'
var year = 1989

var params = {
  TableName: table,
  Key: {
    artist: artist,
    year: year
  }
}

docClient.get(params, function(err, data) {
  if (err) {
    console.log('Unable to read item: ' + JSON.stringify(err))
  } else {
    console.log('GetItem succeeded: ' + JSON.stringify(data))
  }
})
```

Back

Next

Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

SDK Options

Course Navigation

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Object Persistence Interface

- Do not directly perform data plane operations
- Map complex data types to items in a DynamoDB table
- Create objects that represent tables and indexes
- Define the relationships between objects in your program and the tables that store those objects
- Call simple object methods, such as save, load, or delete
- Available in the AWS SDKs for Java and .NET

Back

Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

DynamoDB Local

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

Putitem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

- Downloadable version of DynamoDB
- Test applications locally without accessing the DynamoDB web service
- Docker image or Java .jar file

Java

Download DynamoDB local from links in the documentation

```
java -Djava.library.path=./DynamoDBLocal_lib \
      -jar DynamoDBLocal.jar
```

Docker

All DynamoDB local dependencies and necessary configuration built in

```
docker run -d -p 8000:8000 amazon/dynamodb-local
```

Using the Local Endpoint

AWS CLI

```
aws dynamodb list-tables \
      --endpoint-url http://localhost:8000
```

Python (Boto3)

```
import boto3

client = boto3.client('dynamodb',
                      endpoint_url='http://localhost:8000/')
client.list_tables()
```



DynamoDB Fundamentals

Monitoring and Metrics

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

Putitem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Select the items below regarding monitoring and metrics:

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

[Previous Section](#)

[Next](#)



DynamoDB Fundamentals

Console Metrics and CloudWatch

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

[DynamoDB Architecture](#)[DynamoDB Items](#)

Accessing DynamoDB

[Using the AWS Management Console](#)[Using the AWS CLI](#)[SDK Options](#)[DynamoDB Local](#)

Monitoring and Metrics

Console Metrics and CloudWatch

[Alerts and Alarms](#)[Errors and Codes](#)

Working with Items

[Partitions, Partition and Sort Keys](#)[Performance Units: RCU/WCU](#)[Consistency Model \(Strongly vs. Eventual\)](#)[Scans and Queries](#)[PutItem](#)[Batch Operations](#)

Table Performance

[Provisioned vs. On-Demand Capacity Modes](#)[Auto Scaling](#)

Data Model - Version 1

[Version 1 Data Model Overview](#)

CloudWatch monitors your AWS resources in real time, providing visibility into resource utilization, application performance, and operational health.

- Track metrics (data points over time)
- Create dashboards
- Create alarms
- Create rules for events
- View logs

Helps answer questions like:

- How can I determine how much of my provisioned throughput is being used?
- How can I determine which requests exceed the provisioned throughput limits of a table?
- How can I determine if any system errors occurred?

DynamoDB Metrics

- ConsumedReadCapacityUnits
- ConsumedWriteCapacityUnits
- ProvisionedReadCapacityUnits
- ProvisionedWriteCapacityUnits
- ReadThrottleEvents
- SuccessfulRequestLatency
- SystemErrors
- ThrottledRequests
- UserErrors
- WriteThrottleEvents



DynamoDB Fundamentals

Alerts and Alarms

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

[DynamoDB Architecture](#)[DynamoDB Items](#)

Accessing DynamoDB

[Using the AWS Management Console](#)[Using the AWS CLI](#)[SDK Options](#)[DynamoDB Local](#)

Monitoring and Metrics

[Console Metrics and CloudWatch](#)[Alerts and Alarms](#)[Errors and Codes](#)

Working with Items

[Partitions, Partition and Sort Keys](#)[Performance Units: RCU/WCU](#)[Consistency Model \(Strongly vs. Eventual\)](#)[Scans and Queries](#)[PutItem](#)[Batch Operations](#)

Table Performance

[Provisioned vs. On-Demand Capacity Modes](#)[Auto Scaling](#)

Data Model - Version 1

[Version 1 Data Model Overview](#)

Alarms can be created on metrics, taking an action if the alarm is triggered.

Alarms have three states:

- **INSUFFICIENT:** Not enough data to judge the state — alarms often start in this state.
- **ALARM:** The alarm threshold has been breached (e.g., > 90% CPU).
- **OK:** The threshold has not been breached.

Alarms have a number of key components:

- **Metric:** The data points over time being measured
- **Threshold:** Exceeding this is bad (static or anomaly)
- **Period:** How long the threshold should be bad before an alarm is generated
- **Action:** What to do when an alarm triggers
 - SNS
 - Auto Scaling
 - EC2

[Back](#)

DynamoDB Fundamentals

Errors and Codes

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

[DynamoDB Architecture](#)[DynamoDB Items](#)

Accessing DynamoDB

[Using the AWS Management Console](#)[Using the AWS CLI](#)[SDK Options](#)[DynamoDB Local](#)

Monitoring and Metrics

[Console Metrics and CloudWatch](#)[Alerts and Alarms](#)[Errors and Codes](#)

Working with Items

[Partitions, Partition and Sort Keys](#)[Performance Units: RCU/WCU](#)[Consistency Model \(Strongly vs. Eventual\)](#)[Scans and Queries](#)[PutItem](#)[Batch Operations](#)

Table Performance

[Provisioned vs. On-Demand Capacity Modes](#)[Auto Scaling](#)

Data Model - Version 1

[Version 1 Data Model Overview](#)

Successful DynamoDB requests return the HTTP status code for success (200 OK), along with results from the requested operation.

If the request fails, DynamoDB returns an error with three components:

- HTTP status code (e.g., 400)
- Exception name (e.g., `ResourceNotFoundException`)
- Error message (e.g., `Requested resource not found: Table: tablename not found`)

Most AWS SDKs propagate these low-level errors for you. Handle using try-catch blocks or if-then statements.

Error Retries and Exponential Backoff

AWS SDKs all implement retry logic automatically.

Exponential backoff uses increasing wait times for consecutive error messages. This algorithm helps mitigate request rate errors.

Batch Operations

- Batch operations tolerate the failure of individual requests.
- The most common cause of errors is **throttling**.
- Unprocessed items should be retried using exponential backoff.

[AccessDeniedException](#)[ConditionalCheckFailedException](#)[IncompleteSignatureException](#)[ItemCollectionSizeLimitExceededException](#)[LimitExceededException](#)[MissingAuthenticationTokenException](#)[ProvisionedThroughputExceededException](#)[RequestLimitExceeded](#)[ResourceInUseException](#)[ResourceNotFoundException](#)[ThrottlingException](#)[UnrecognizedClientException](#)[ValidationException](#)[InternalServerError](#)[ServiceUnavailable](#)

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Select the items below regarding working with items:

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Previous Section

Next

DynamoDB Fundamentals

Partitions, Partition and Sort Keys

Course Navigation

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

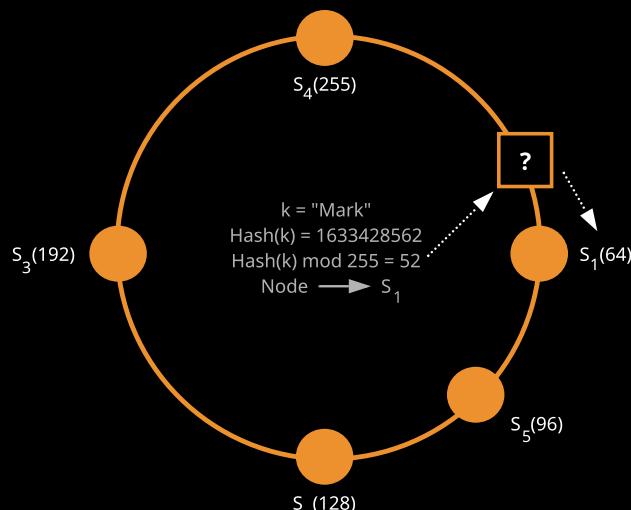
- The partition key or partition and sort key are used to uniquely **write** or **read** an item.
- Partition key is used for partition selection via the DynamoDB internal hash function
- One partition holds **10 GB** of data and supports up to **3,000** read capacity units (RCU) or **1,000** write capacity units (WCU).
- Partition key is used to select the item(s) to read/write.
- Sort key is used to range select (e.g., `begins_with`) or to order results.
- Sort keys may not be used on their own.

Consistent Hashing

Key: k

N storage nodes

Selected node: $\text{Hash}(k) \bmod N$



[Back to Main](#)

[Section 5](#)



Linux Academy

DynamoDB Fundamentals

Performance Units: RCU/WCU

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Provisioned Throughput

Maximum amount of capacity an application can consume from a table or index.

Throttled requests: ProvisionedThroughputExceededException

Eventually vs. Strongly Consistent Read

Eventually consistent reads might include stale data.

Strongly consistent reads are always up to date but are subject to network delays.

Read Capacity Units (RCUs)

One RCU represents one strongly consistent read request per second, or two eventually consistent read requests, for an item up to 4 KB in size. Transactional read requests require 2 RCUs for items up to 4 KB.

Filtered query or scan results consume full read capacity.

For an **8 KB** item size:

- 2 RCUs for one strongly consistent read
- 1 RCU for an eventually consistent read
- 4 RCUs for a transactional read

Write vs. Transactional Write

Writes are eventually consistent within one second or less.

Write Capacity Units (WCUs)

One WCU represents one write per second for an item up to 1 KB in size.

Transactional write requests require 2 WCUs for items up to 1 KB.

For a **3 KB** item size:

- Standard: 3 WCUs
- Transactional: 6 WCUs

Provisioned Throughput Calculations



DynamoDB Fundamentals

Performance Units: RCU/WCU

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Provisioned Throughput

Maximum amount of capacity an application can consume from a table or index.

Provisioned Throughput Calculations



A system needs to store 60 patient records of 1.5 KB, each, every minute. What WCU should you allocate on the patient record table?

- 60 records per minute = ~1 per second (and the DDB RCU/WCU buffer can smooth this out if not)
- Each record is 1.5 KB. 1 WCU = 1 KB per second, so each record requires 2 WCU.
- A WCU setting of 2 is required on the table.

A weather application reads data from a DynamoDB table. Each item in the table is 7 KB in size. How many RCUs should be set on the table to allow for 10 reads per second?

- 1 item is 7 KB, which is 2 RCU (1 RCU is 4 KB).
- 10 reads per second for 7 KB items = 20 RCU
- But the question didn't specify if eventual or strong consistency is required. The default is eventual, which allows for 2 reads of 4 KB per second for 1 RCU.
- Assuming eventually consistent reads, the answer is 10 RCU.

- Standard: 3 WCUs
- Transactional: 6 WCUs

Provisioned Throughput Calculations



DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

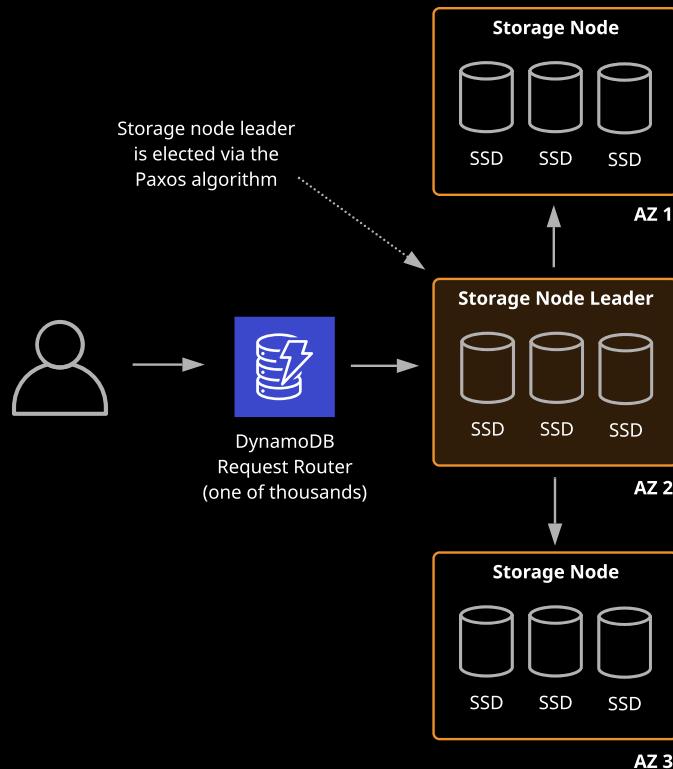
Version 1 Data Model Overview

[Back to Main](#)

DynamoDB Fundamentals

Consistency Model (Strongly vs. Eventual)

PutItem Request

[Next](#)[Section 5](#)

Linux Academy

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

[DynamoDB Architecture](#)[DynamoDB Items](#)

Accessing DynamoDB

[Using the AWS Management Console](#)[Using the AWS CLI](#)[SDK Options](#)[DynamoDB Local](#)

Monitoring and Metrics

[Console Metrics and CloudWatch](#)[Alerts and Alarms](#)[Errors and Codes](#)

Working with Items

[Partitions, Partition and Sort Keys](#)[Performance Units: RCU/WCU](#)[Consistency Model \(Strongly vs. Eventual\)](#)[Scans and Queries](#)[PutItem](#)[Batch Operations](#)

Table Performance

[Provisioned vs. On-Demand Capacity Modes](#)[Auto Scaling](#)

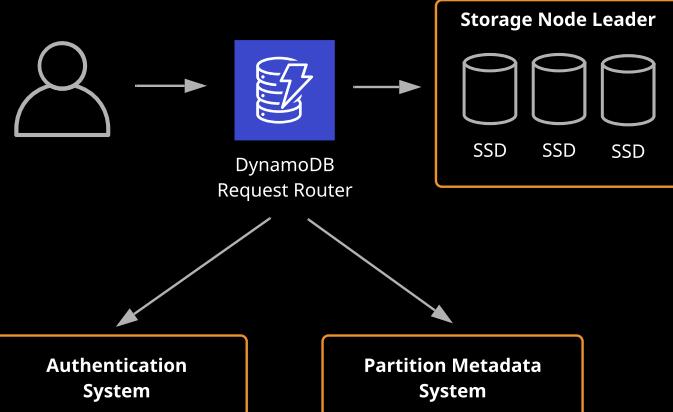
Data Model - Version 1

[Version 1 Data Model Overview](#)

DynamoDB Fundamentals

Consistency Model (Strongly vs. Eventual)

GetItem Request

[Back](#)[Next](#)[Back to Main](#)[Section 5](#)

Linux Academy

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

[DynamoDB Architecture](#)[DynamoDB Items](#)

Accessing DynamoDB

[Using the AWS Management Console](#)[Using the AWS CLI](#)[SDK Options](#)[DynamoDB Local](#)

Monitoring and Metrics

[Console Metrics and CloudWatch](#)[Alerts and Alarms](#)[Errors and Codes](#)

Working with Items

[Partitions, Partition and Sort Keys](#)[Performance Units: RCU/WCU](#)[Consistency Model \(Strongly vs. Eventual\)](#)[Scans and Queries](#)[PutItem](#)[Batch Operations](#)

Table Performance

[Provisioned vs. On-Demand Capacity Modes](#)[Auto Scaling](#)

Data Model - Version 1

[Version 1 Data Model Overview](#)

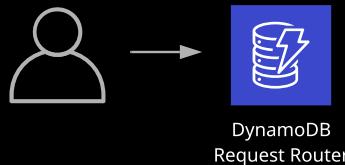
DynamoDB Fundamentals

Consistency Model (Strongly vs. Eventual)

Eventual Consistency - GetItem

Request router randomly selects a storage node for the partition

Storage Node



Storage Node Leader



Storage Node



- 1 in 3 chance that you've requested stale data (worst case).
- Eventually consistent reads are 50% the "cost" of strongly consistent.
- Strongly consistent reads require 2 of 3 matches across replicas.

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Scan

- Returns all items and attributes for a given table
- Filtering results do not reduce RCU consumption; they simply discard data
- Eventually consistent by default, but the ConsistentRead parameter can enable strongly consistent scans
- Limit the number of items returned
- A single query returns results that fit within 1 MB
- Pagination can be used to retrieve more than 1 MB
- Parallel scans can be used to improve performance
- Prefer query over scan when possible; occasional real-world use is okay
- If you are repeatedly using scans to filter on the same non-PK/SK attribute, consider creating a secondary index

Query

- Find items based on primary key values
- Query limited to PK, PK+SK, or secondary indexes
- Requires PK attribute
- Returns all items with that PK value
- Optional SK attribute and comparison operator to refine results
- Filtering results do not reduce RCU consumption; they simply discard data
- Eventually consistent by default, but the ConsistentRead parameter can enable strongly consistent queries
- Querying a partition only scans that one partition
- Limit the number of items returned
- A single query returns results that fit within 1 MB
- Pagination can be used to retrieve more than 1 MB



DynamoDB Fundamentals

PutItem

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

PutItem

- Creates a new item
- Replaces existing item with the same key

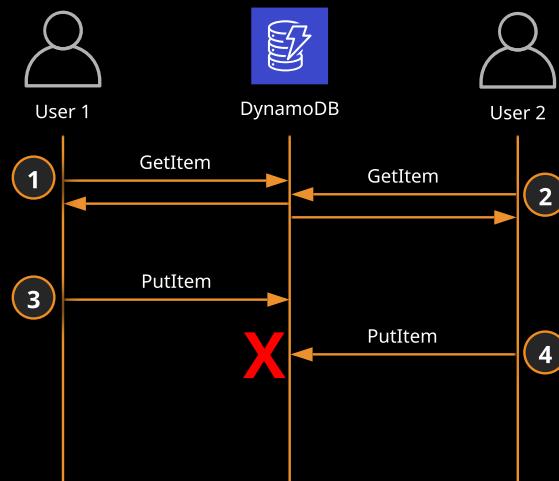
UpdateItem

- Creates a new item if the specified key does not exist; otherwise, it modifies an existing item's attributes

DeleteItem

- Deletes the item with the specified key

Write Conflicts



Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

BatchGetItem

- Returns attributes for multiple items from multiple tables
- Request using primary key
- Returns up to 16 MB of data, up to 100 items
- Get unprocessed items exceeding limits via UnprocessedKeys
- Eventually consistent by default, but the ConsistentRead parameter can enable strongly consistent scans
- Retrieves items in parallel to minimize latency

BatchWriteItem

- Puts or deletes multiple items in multiple tables
- Writes up to 16 MB of data, up to 25 put or delete requests
- Get unprocessed items exceeding limits via UnprocessedItems
- Conditions are not supported for performance reasons
- Threading may be used to write items in parallel



DynamoDB Fundamentals

Table Performance

Course Navigation

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Select the items below regarding table performance:

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Previous Section

Next

Back to Main

Section 5



Linux Academy

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

[DynamoDB Architecture](#)[DynamoDB Items](#)

Accessing DynamoDB

[Using the AWS Management Console](#)[Using the AWS CLI](#)[SDK Options](#)[DynamoDB Local](#)

Monitoring and Metrics

[Console Metrics and CloudWatch](#)[Alerts and Alarms](#)[Errors and Codes](#)

Working with Items

[Partitions, Partition and Sort Keys](#)[Performance Units: RCU/WCU](#)[Consistency Model \(Strongly vs. Eventual\)](#)[Scans and Queries](#)[PutItem](#)[Batch Operations](#)

Table Performance

[Provisioned vs. On-Demand Capacity Modes](#)[Auto Scaling](#)

Data Model - Version 1

[Version 1 Data Model Overview](#)

DynamoDB Fundamentals

Provisioned vs. On-Demand Capacity Modes

Provisioned Capacity

- Minimum capacity required
- Able to set a budget (maximum capacity)
- Subject to throttling
- Auto scaling available
- Risk of underprovisioning — monitor your metrics
- Lower price per API call
- \$0.00065 per WCU-hour (us-east-1)
- \$0.00013 per RCU-hour (us-east-1)
- \$0.25 per GB-month (first 25 GB is free)
- Example:
 - 25 WCU, 25 RCU, 25 GB
 - 25 WCU x \$0.00065 per hour x 24 hours x 30 days = \$11.70
 - 25 RCU x \$0.00013 per hour x 24 hours x 30 days = \$2.34
 - 25 GB = \$0.00
 - Total Bill: \$14.04

On-Demand Capacity

- No minimum capacity; pay more per request than provisioned capacity
- Idle tables not charged for read/write, but only for storage and backups
- No capacity planning required — just make API calls
- Eliminates the tradeoffs of over- or under-provisioning
- Use on-demand for new product launches
- Switch to provisioned once a steady state is reached
- \$1.25 per million WCU (us-east-1)
- \$0.25 per million RCU (us-east-1)
- Example:
 - You host a web service where you charge per API call
 - On-demand prevents being throttled because you never want to be down
 - 4,000,000 WCU per month
 - 8,000,000 RCU per month
 - 30 GB storage
 - \$1.25 per million WCU x 4 million WCU = \$5.00
 - \$0.25 per million RCU x 8 million RCU = \$2.00
 - 10 GB x \$0.25 = \$2.50
 - Total Bill: \$9.50



DynamoDB Fundamentals

Auto Scaling

Course Navigation

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

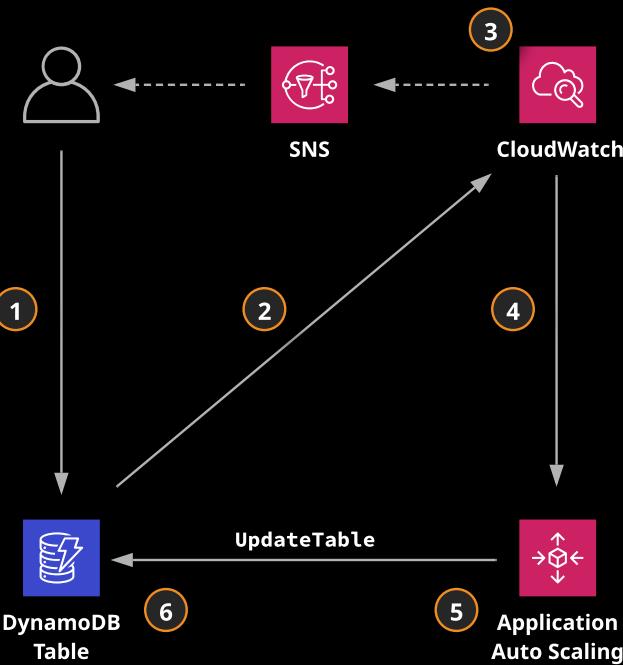
Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

- Auto scaling is best suited for predictable, gradually changing traffic patterns.
- You may still manually change provisioned throughput settings.
- Recommended to apply auto scaling uniformly to tables and global secondary indexes (GSIs).



[Back to Main](#)

[Section 5](#)



Linux Academy

DynamoDB Fundamentals

Version 1 Data Model Overview

DynamoDB Fundamentals

Section 4

Tables and Item Architecture

DynamoDB Architecture

DynamoDB Items

Accessing DynamoDB

Using the AWS Management Console

Using the AWS CLI

SDK Options

DynamoDB Local

Monitoring and Metrics

Console Metrics and CloudWatch

Alerts and Alarms

Errors and Codes

Working with Items

Partitions, Partition and Sort Keys

Performance Units: RCU/WCU

Consistency Model (Strongly vs. Eventual)

Scans and Queries

PutItem

Batch Operations

Table Performance

Provisioned vs. On-Demand Capacity Modes

Auto Scaling

Data Model - Version 1

Version 1 Data Model Overview

Table Name: album**PK****id****Secondary Item Attributes**

album_art

artist_id

format

price

...

Partition Key: id (N)**Secondary Attributes:** album_art (S), artist_id (N), format (S), price (N), title (S), year (N)**Table Name:** artist**PK****id**

name

Partition Key: id (N)**Secondary Attributes:** name (S)**Table Name:** track**PK****id**

album_id

name

number

length

Partition Key: id (N)**Secondary Attributes:** album_id (N), name (S), number (N), length (N)**Table Name:** user**PK****email**

password

Partition Key: email (S)**Secondary Attributes:** password (S)

Intermediate DynamoDB

Table Indexes

Course Navigation

Intermediate DynamoDB

Section 5

Table Indexes

Indexes: Part 1

Indexes: Part 2

Importing Tables Using AWS Database Migration Service

Backup and Recovery

On-Demand vs. Continuous (PITR) Backup

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

Hot and Cold Partition Imbalance

Data Consistency and Management

Conditional and Update Expressions

Transactions

Time to Live (TTL)

Select the items below regarding table indexes:

Indexes: Part 1

Indexes: Part 2

Importing Tables Using AWS Database Migration Service

[Next Topics](#)

[Back to Main](#)



Linux Academy

Intermediate DynamoDB

Section 5

Table Indexes

Indexes: Part 1

Indexes: Part 2

Importing Tables Using AWS Database Migration Service

Backup and Recovery

On-Demand vs. Continuous (PITR) Backup

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

Hot and Cold Partition Imbalance

Data Consistency and Management

Conditional and Update Expressions

Transactions

Time to Live (TTL)

Next Topics

Back to Main

Secondary Indexes

Data structure that contains a **subset of attributes** from a table, along with an **alternate key** to support query operations. You can **retrieve data from the index using a query**, just like with a table. A table can have multiple secondary indexes.

Global Secondary Index: Index with a partition key and a sort key that can be different from those on the base table. The primary key of a GSI can be either simple (partition key) or composite (partition key and sort key). Strong consistency is unsupported.

Local Secondary Index: Index that has the same partition key as the base table but a different sort key. The primary key of a LSI must be composite (partition and sort key). Must be created at the time of table creation.

Partition Key		Sort Key	
username	email	order_id	total
mark	mark@linuxacademy.com	223	4096.64
terry	tcox@linuxacademy.com	224	1024.16

Order Table

Partition Key		Sort Key	
username	email	order_id	total
mark	mark@linuxacademy.com	223	4096.64
terry	tcox@linuxacademy.com	224	1024.16

Global Secondary Index (GSI)

Partition Key		Sort Key	
username	email	order_id	total
mark	mark@linuxacademy.com	223	4096.64
terry	tcox@linuxacademy.com	224	1024.16

Local Secondary Index (LSI)



Intermediate DynamoDB

Indexes: Part 2

Course Navigation

Intermediate DynamoDB

Section 5

Table Indexes

[Indexes: Part 1](#)

[Indexes: Part 2](#)

Importing Tables Using AWS Database Migration Service

Backup and Recovery

On-Demand vs. Continuous (PITR) Backup

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

Hot and Cold Partition Imbalance

Data Consistency and Management

Conditional and Update Expressions

Transactions

Time to Live (TTL)

[Next Topics](#)

[Back to Main](#)

Index Queries and Projection Expressions

A projection is the set of attributes copied from a table into a secondary index. The partition key and sort key of the table are always projected into the index.

ALL — All of the table attributes are projected into the index.

KEYS_ONLY — Only the index and primary keys are projected into the index.

INCLUDE — Only the specified table attributes are projected into the index.

Using **GetItem**, **Query**, or **Scan** operations return all attributes by default. Use **Projection Expressions** to get a subset of these attributes.

Example: Find all albums by artist_id, returning only title and price:

```
response = table.query(  
    IndexName="artist_id-index",  
    KeyConditionExpression=Key("artist_id").eq(192),  
    ProjectionExpression="title, price",  
)
```

Response:

```
{"price": "30.94", "title": "The Game"}  
{"price": "7.45", "title": "A Night at the Opera"}  
{"price": "7.18", "title": "Crazy Little Thing Called Love"}  
{"price": "7.22", "title": "The Invisible Man"}  
{"price": "48.97", "title": "Don't Stop Me Now"}  
{"price": "3.63", "title": "Keep Yourself Alive"}  
{"price": "35.63", "title": "A Day at the Races"}  
{"price": "35.31", "title": "Under Pressure"}  
{"price": "11.73", "title": "I Want to Break Free"}  
...  
{"price": "22.6", "title": "Bohemian Rhapsody"}
```

[Next](#)



Linux Academy

Intermediate DynamoDB

Section 5

Table Indexes

[Indexes: Part 1](#)

[Indexes: Part 2](#)

Importing Tables Using AWS Database Migration Service

Backup and Recovery

On-Demand vs. Continuous (PITR) Backup

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

Hot and Cold Partition Imbalance

Data Consistency and Management

Conditional and Update Expressions

Transactions

Time to Live (TTL)

Sparse Indexes

DynamoDB writes a corresponding index item only if the index sort key value is present in the item. If the sort key doesn't appear in every table item, the index is said to be **sparse**.

In other words, the index is "sparse" because it does not include all items of the parent table.

Global secondary indexes are sparse by default. When you create a GSI, you specify a partition key and optionally a sort key. Only items in the parent table that contain those attributes appear in the index.

By creating sparse indexes, you can provision GSIs with lower write throughput than that of the parent table. This can possibly save a lot of money that would otherwise be spent on storage (for no additional benefits).

[Back](#)

[Next Topics](#)

[Back to Main](#)



Linux Academy

Intermediate DynamoDB

Importing Tables Using AWS Database Migration Service

Course Navigation

Intermediate DynamoDB

Section 5

Table Indexes

Indexes: Part 1

Indexes: Part 2

Importing Tables Using AWS Database Migration Service

Backup and Recovery

On-Demand vs. Continuous (PITR) Backup

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

Hot and Cold Partition Imbalance

Data Consistency and Management

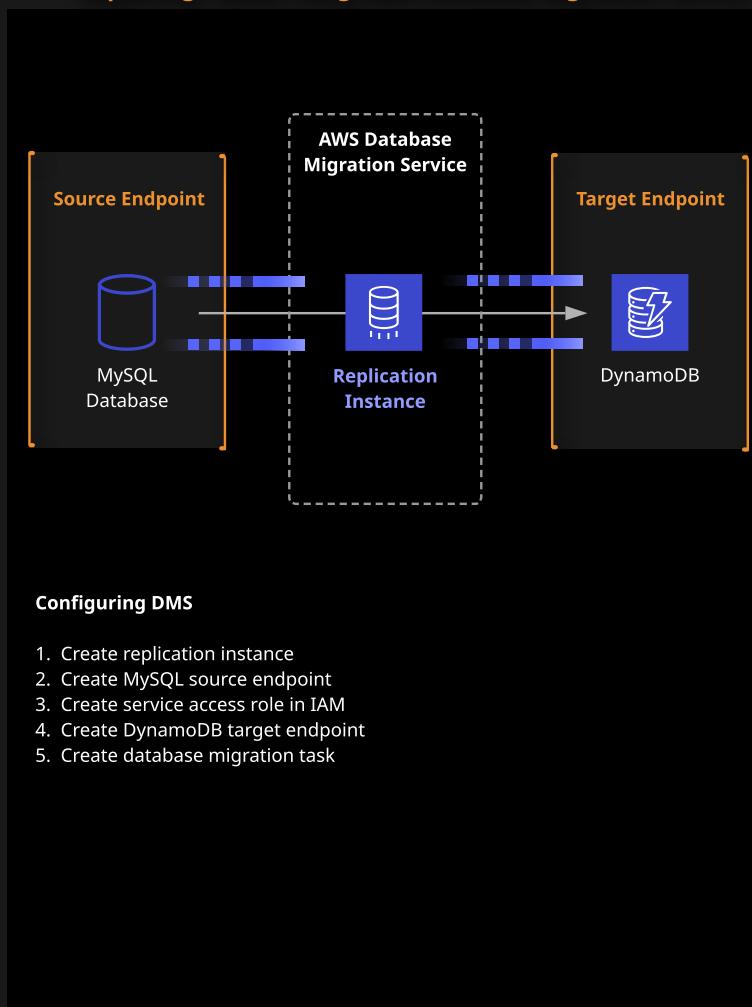
Conditional and Update Expressions

Transactions

Time to Live (TTL)

Next Topics

Back to Main



Configuring DMS

1. Create replication instance
2. Create MySQL source endpoint
3. Create service access role in IAM
4. Create DynamoDB target endpoint
5. Create database migration task



Linux Academy

Intermediate DynamoDB

Section 5

Table Indexes

Indexes: Part 1

Indexes: Part 2

Importing Tables Using AWS Database Migration Service

Backup and Recovery

On-Demand vs. Continuous (PITR) Backup

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

Hot and Cold Partition Imbalance

Data Consistency and Management

Conditional and Update Expressions

Transactions

Time to Live (TTL)

[Next Topics](#)

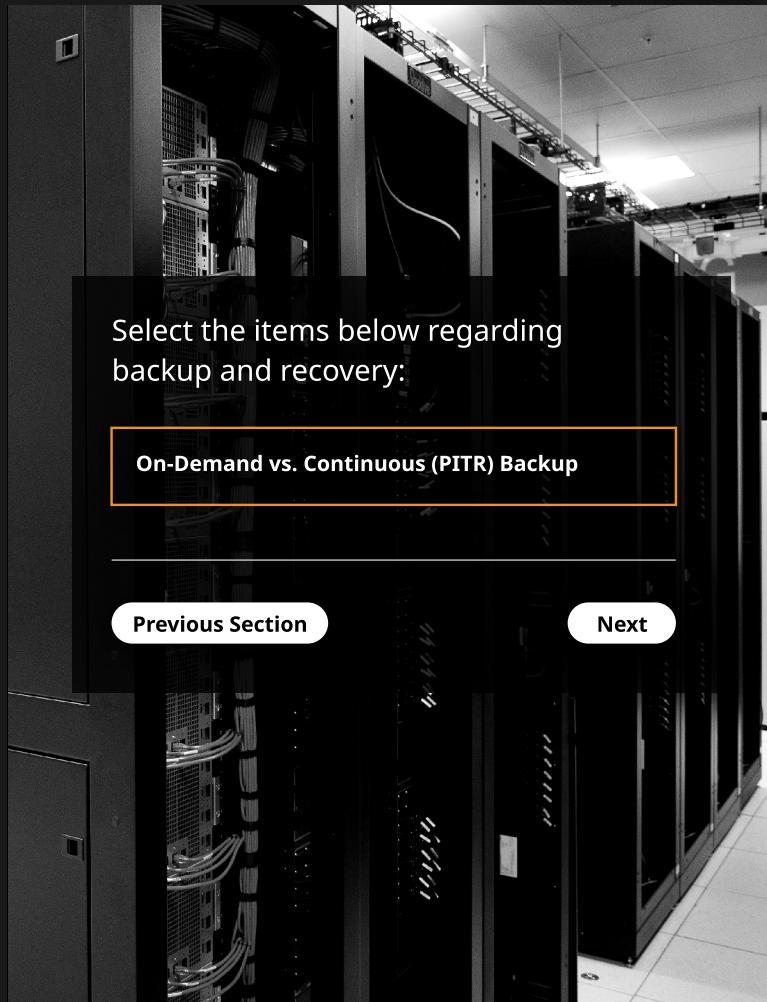
[Back to Main](#)

Select the items below regarding backup and recovery:

On-Demand vs. Continuous (PITR) Backup

[Previous Section](#)

[Next](#)



Intermediate DynamoDB

Section 5

Table Indexes

[Indexes: Part 1](#)

[Indexes: Part 2](#)

[Importing Tables Using AWS Database Migration Service](#)

Backup and Recovery

[On-Demand vs. Continuous \(PITR\) Backup](#)

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

Hot and Cold Partition Imbalance

Data Consistency and Management

Conditional and Update Expressions

Transactions

Time to Live (TTL)

[Next Topics](#)

[Back to Main](#)

Intermediate DynamoDB

On-Demand vs. Continuous (PITR) Backup

On-Demand Backup and Restore

Create full backups of your tables at any time in either the AWS Management Console or with a single API call.

- Backup and restore actions execute with zero impact on table performance or availability.
- Backups are consistent within seconds and retained until deleted.
- Backup and restore operates within the same region as the source table.

Point-in-Time Recovery (PITR)

Helps protect your DynamoDB tables from accidental writes or deletes. You can restore your data to any point in time in the last **35 days**.

- DynamoDB maintains **incremental** backups of your data.
- Point-in-time recovery is not enabled by default.
- The latest restorable timestamp is typically five minutes in the past.

After restoring a table, you must manually set up the following on the restored table:

- Auto scaling policies
- AWS Identity and Access Management (IAM) policies
- Amazon CloudWatch metrics and alarms
- Tags
- Stream settings
- Time to Live (TTL) settings
- Point-in-time recovery settings



Linux Academy

Intermediate DynamoDB

Advanced Performance and Scaling Considerations

Intermediate DynamoDB

Section 5

Table Indexes

Indexes: Part 1

Indexes: Part 2

Importing Tables Using AWS Database Migration Service

Backup and Recovery

On-Demand vs. Continuous (PTTR) Backup

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

Hot and Cold Partition Imbalance

Data Consistency and Management

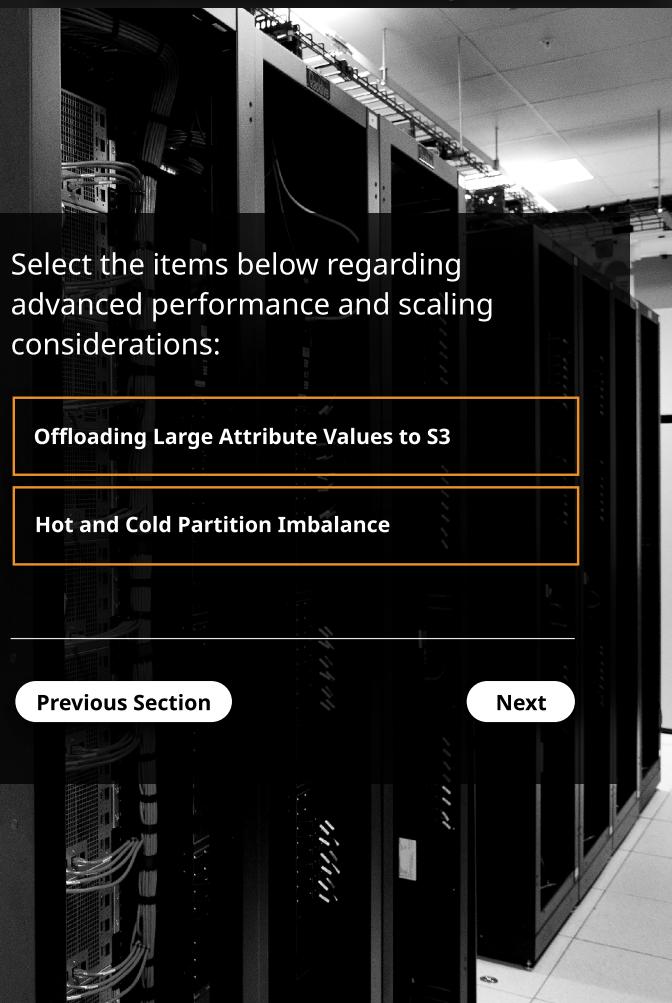
Conditional and Update Expressions

Transactions

Time to Live (TTL)

[Next Topics](#)

[Back to Main](#)



Select the items below regarding advanced performance and scaling considerations:

Offloading Large Attribute Values to S3

Hot and Cold Partition Imbalance

[Next](#)

[Previous Section](#)



Linux Academy

Intermediate DynamoDB

Offloading Large Attribute Values to S3

Intermediate DynamoDB

Section 5

Table Indexes

Indexes: Part 1

Indexes: Part 2

Importing Tables Using AWS Database Migration Service

Backup and Recovery

On-Demand vs. Continuous (PITR) Backup

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

Hot and Cold Partition Imbalance

Data Consistency and Management

Conditional and Update Expressions

Transactions

Time to Live (TTL)

Large Attribute Values in Attributes

Table Name: album

PK

id

Secondary Item Attributes

album_art

artist_id

format

price

...

350 KB

1 KB

Offloading Large Attribute Values to S3

Table Name: album

PK

id

Secondary Item Attributes

album_art

artist_id

format

price

...

S3 URI



S3 Bucket

Next Topics

Back to Main



Linux Academy

Intermediate DynamoDB

Hot and Cold Partition Imbalance

Intermediate DynamoDB

Section 5

Table Indexes

Indexes: Part 1

Indexes: Part 2

Importing Tables Using AWS Database Migration Service

Backup and Recovery

On-Demand vs. Continuous (PTTR) Backup

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

Hot and Cold Partition Imbalance

Data Consistency and Management

Conditional and Update Expressions

Transactions

Time to Live (TTL)

[Next Topics](#)

[Back to Main](#)

Resharding

$$(RCU/3,000) + (WCU/1,000) = \# \text{ Partitions}$$

1



2



3



Adaptive Capacity

Example Table

Total provisioned capacity = 400 WCU

Total consumed capacity = 300 WCU



Intermediate DynamoDB

Section 5

Table Indexes

Indexes: Part 1

Indexes: Part 2

Importing Tables Using AWS Database Migration Service

Backup and Recovery

On-Demand vs. Continuous (PITR) Backup

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

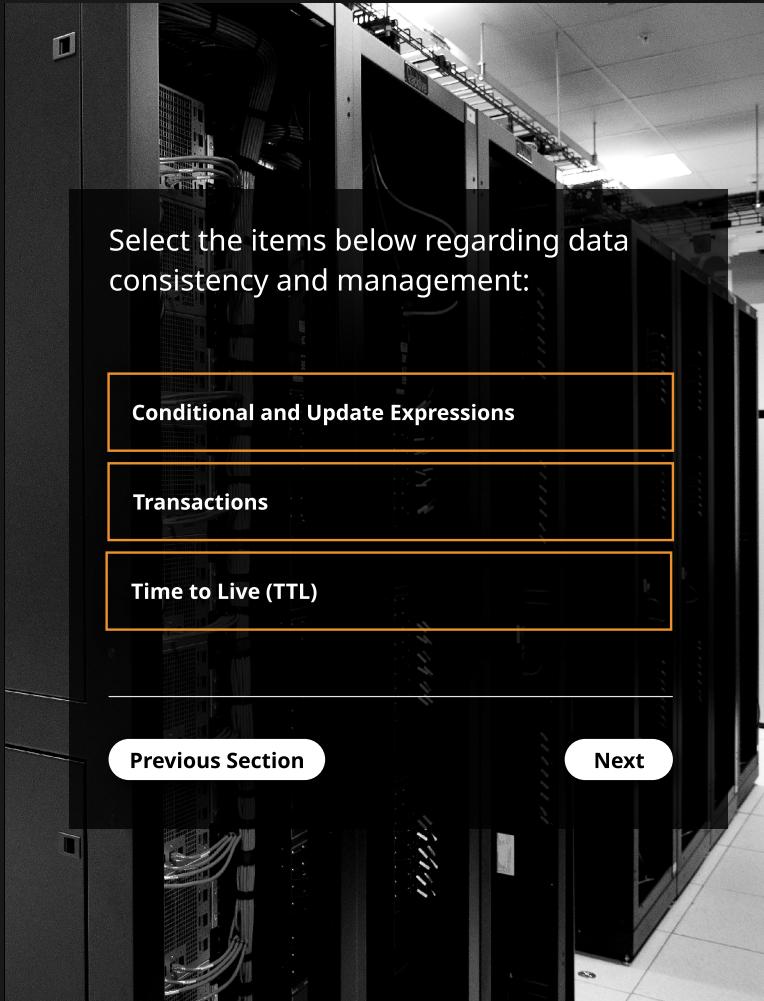
Hot and Cold Partition Imbalance

Data Consistency and Management

Conditional and Update Expressions

Transactions

Time to Live (TTL)



Select the items below regarding data consistency and management:

Conditional and Update Expressions

Transactions

Time to Live (TTL)

Next Topics

Previous Section

Next

Back to Main



Linux Academy

Intermediate DynamoDB

Conditional and Update Expressions

Intermediate DynamoDB

Section 5

Table Indexes

Indexes: Part 1

Indexes: Part 2

Importing Tables Using AWS Database Migration Service

Backup and Recovery

On-Demand vs. Continuous (PTTR) Backup

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

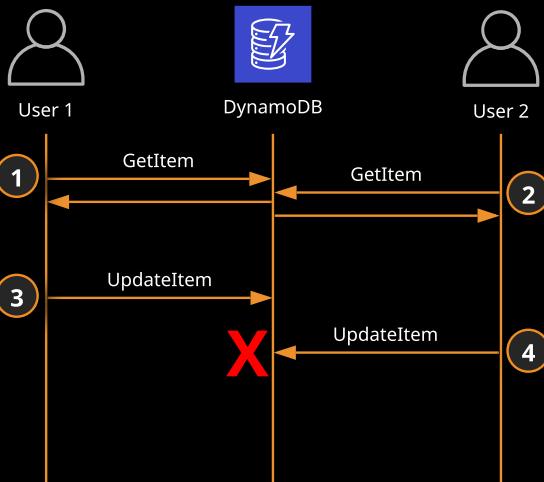
Hot and Cold Partition Imbalance

Data Consistency and Management

Conditional and Update Expressions

Transactions

Time to Live (TTL)



Conditional Updates

Performs a put/update/delete operation only if a condition is met:

- Check attribute value
- Check if attribute exists

Atomic counter can be used to increment or decrement the value of an existing attribute without interfering with other write requests.

[Next Topics](#)

[Back to Main](#)



Linux Academy

Intermediate DynamoDB

Transactions

Course Navigation

Intermediate DynamoDB

Section 5

Table Indexes

Indexes: Part 1

Indexes: Part 2

Importing Tables Using AWS Database Migration Service

Backup and Recovery

On-Demand vs. Continuous (PITR) Backup

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

Hot and Cold Partition Imbalance

Data Consistency and Management

Conditional and Update Expressions

Transactions

Time to Live (TTL)

Next Topics

Back to Main

Transactions enable developers to easily perform atomic writes and isolated reads across multiple items and tables.

TransactWriteItems:

- Transact up to 25 items or 4 MB
- Evaluate conditions
- If all conditions are simultaneously true, perform write operations

TransactGetItems:

- Transact up to 25 items or 4 MB
- Return a consistent, isolated snapshot of all items

DynamoDB performs **two underlying reads or writes** of every item in the transaction, one to prepare the transaction and one to commit the transaction.

Apply transactions to items in:

- Same partition key or across partition keys
- Same table or across tables
- Same region only
- Same account only
- DynamoDB tables only

Transactions scale like the rest of DynamoDB:

- Unlimited concurrent transactions
- Low latency, high availability
- Scales horizontally
- Up to you to design to avoid hot keys

Reasons for transaction failures:

- Precondition failure
- Insufficient capacity
- Transactional conflicts
- Transaction still in progress
- Service error
- Malformed request
- Permissions



Linux Academy

Intermediate DynamoDB

Section 5

Table Indexes

Indexes: Part 1

Indexes: Part 2

Importing Tables Using AWS Database Migration Service

Backup and Recovery

On-Demand vs. Continuous (PTTR) Backup

Advanced Performance and Scaling Considerations

Offloading Large Attribute Values to S3

Hot and Cold Partition Imbalance

Data Consistency and Management

Conditional and Update Expressions

Transactions

Time to Live (TTL)

Time to Live (TTL) lets you define when table items expire so they can be automatically deleted from the database.

- No cost to use TTL
- Reduce storage costs by limiting storage usage to relevant items only
- Does not consume provisioned throughput

Use cases:

- Session data
- Event logs
- Pattern searching and usage tracking
- Debugging and analytics
- Other temporary sensitive data for contractual or regulatory compliance

How It Works

1. TTL compares the current time to the defined TTL attribute of an item.
2. If current time > item's TTL value, then the item is marked for deletion.
3. DynamoDB typically deletes expired items within **48 hours** of expiration.
4. Items are removed from LSIs and GSIs automatically using an eventually consistent delete operation.



The **epoch** time format is the number of seconds elapsed since 12:00:00 AM January 1, 1970 UTC.

Due to the potential delay between expiration and deletion time, you might get expired items when you query for items. Use a **filter expression** to return only items where the item's TTL expiration > current time.

[Next Topics](#)[Back to Main](#)

Linux Academy

Intermediate DynamoDB

Data Resilience, Security, and Encryption

Course Navigation

Intermediate DynamoDB

Section 5 (Continued)

Data Resilience, Security, and Encryption

Global Tables

Encryption

DynamoDB VPC Endpoints

Database Model v2

Data Model v2 Review

Advanced DynamoDB

Section 6

Select the items below regarding data resilience, security, and encryption:

Global Tables

Encryption

DynamoDB VPC Endpoints

Previous Section

Next

Previous Topics

Back to Main



Linux Academy

Intermediate DynamoDB

Global Tables

Course Navigation

Intermediate DynamoDB

Section 5 (Continued)

Data Resilience, Security, and Encryption

Global Tables

Encryption

DynamoDB VPC Endpoints

Database Model v2

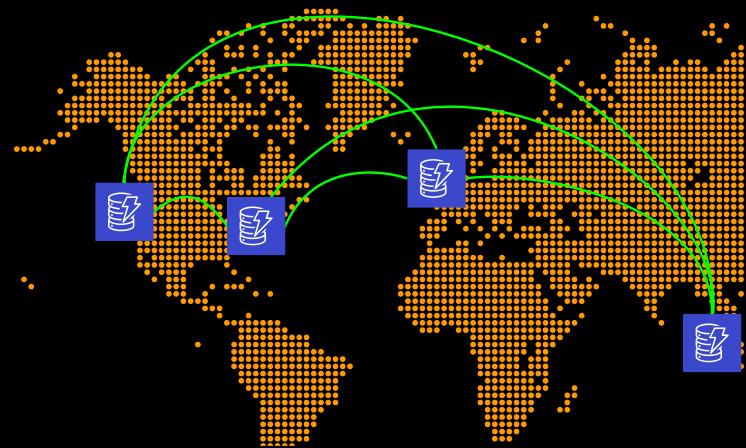
Data Model v2 Review

Advanced DynamoDB

Section 6

Global tables are a fully managed, multi-master, multi-region replication solution.

- Build high performance, globally distributed applications
- Low latency reads and writes to locally available tables
- Multi-region redundancy for disaster recovery (DR) or high availability (HA) applications
- Easy to set up and no application rewrites required
- Replication latency typically under one second



[Previous Topics](#)

[Back to Main](#)



Linux Academy

Intermediate DynamoDB

Section 5 (Continued)

Data Resilience, Security, and Encryption

Global Tables

Encryption

DynamoDB VPC Endpoints

Database Model v2

Data Model v2 Review

Advanced

DynamoDB

Section 6

DynamoDB protects data stored **at rest** and **in transit** between on-premises clients and DynamoDB, and between DynamoDB and other AWS resources within the same AWS Region.

Encryption in transit uses HTTPS.

Encryption at rest uses AWS Key Management Service (KMS) to encrypt all table data, including:

- Primary key
- LSIs and GSIs
- Streams
- Global tables
- Backups
- DAX clusters

When creating a new table, you may choose:

- AWS owned customer master key (CMK)
 - Default option; key owned by DynamoDB; no charges
- AWS managed CMK
 - Key stored in AWS account; managed by AWS KMS
 - View the CMK and its key policy (policy can't be changed)
 - Audit encryption/decryption of your table using CloudTrail

DynamoDB Encryption Client

- Java and Python libraries that encrypt your data **prior** to sending to DynamoDB
- Use cryptographic materials from any source, including your own cryptography service



[Previous Topics](#)

[Back to Main](#)



Linux Academy

Intermediate DynamoDB

DynamoDB VPC Endpoints

Intermediate DynamoDB

Section 5 (Continued)

Data Resilience, Security, and Encryption

Global Tables

Encryption

DynamoDB VPC Endpoints

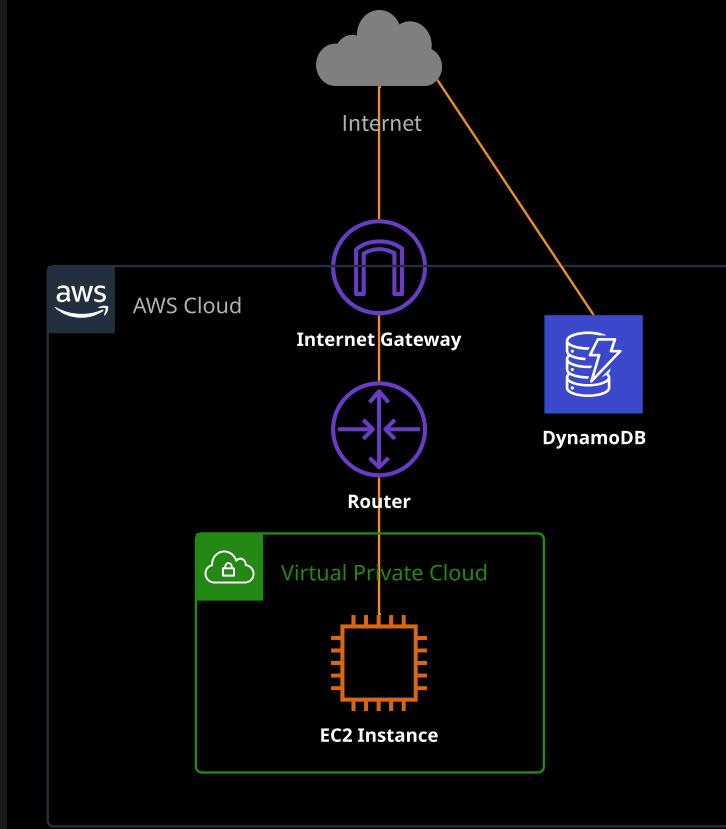
Database Model v2

Data Model v2 Review

Advanced DynamoDB

Section 6

Without VPC Endpoints



Next

Previous Topics

Back to Main



Linux Academy

Intermediate DynamoDB

Section 5 (Continued)

Data Resilience, Security, and Encryption

Global Tables

Encryption

DynamoDB VPC Endpoints

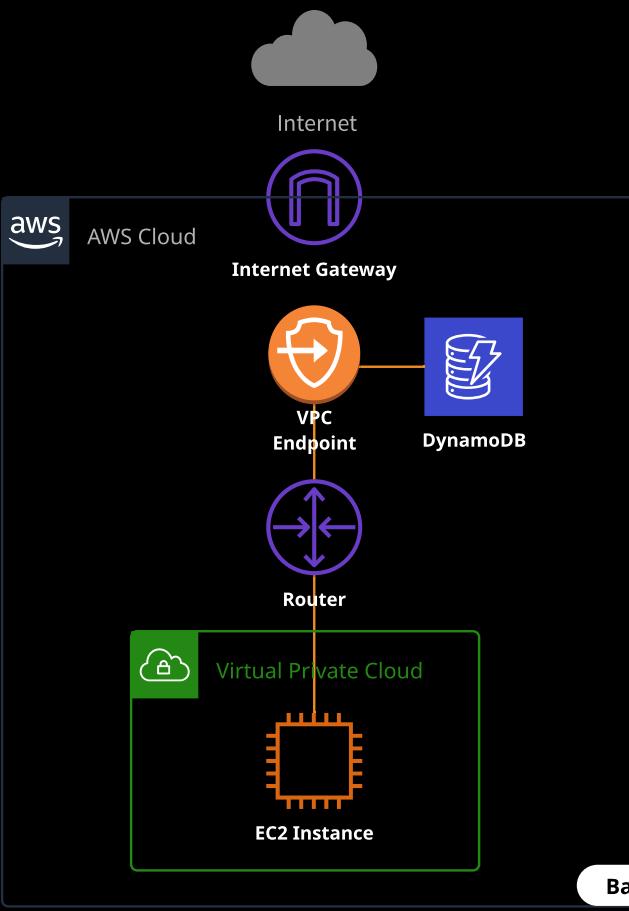
Database Model v2

Data Model v2 Review

Advanced DynamoDB

Section 6

With VPC Endpoints



[Previous Topics](#)

[Back to Main](#)



Linux Academy

Intermediate DynamoDB

Database Model v2

Course Navigation

Intermediate DynamoDB

Section 5 (Continued)

Data Resilience, Security, and Encryption

Global Tables

Encryption

DynamoDB VPC Endpoints

Database Model v2

Data Model v2 Review

Advanced DynamoDB

Section 6

Select the item below regarding database model v2:

Data Model v2 Review

Previous Section

Next

Previous Topics

Back to Main



Linux Academy

Intermediate DynamoDB

Section 5 (Continued)

Data Resilience, Security, and Encryption

Global Tables

Encryption

DynamoDB VPC Endpoints

Database Model v2

Data Model v2 Review

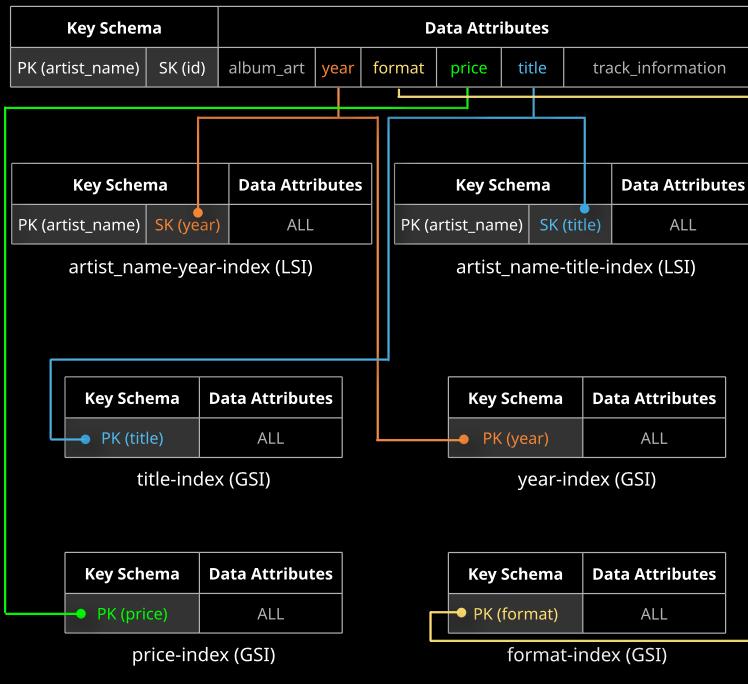
Advanced DynamoDB

Section 6

This version of the Pinehead Records data model features:

- Better table structure (single hierarchical table)
- Local and global secondary indexes
- User accounts in DynamoDB

pinehead_records_s2



[Previous Topics](#)

[Back to Main](#)



Linux Academy

Advanced DynamoDB

Streams and Triggers

Course Navigation

Advanced DynamoDB

Section 6

Streams and Triggers

Streams
Triggers (Lambda Invoke)

Advanced Index Usage
Selective Write Sharding
Aggregation with Streams

Implementing
High-Performance
Architectures

Static Data Dumps
(streams to static file)
DynamoDB Accelerator
(DAX) Architecture
SQS Write Buffer

Advanced Security

Federated and
Fine-Grained Access
Auditing Admin Access
Using CloudTrail

Database Model v3

Database Model v3 Review

Integration

Amazon EMR
Amazon Elasticsearch
Service

Select the items below regarding streams and triggers:

Streams

Triggers (Lambda Invoke)

Previous Section

Next

Back to Main



Linux Academy

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams to static file)

DynamoDB Accelerator (DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and Fine-Grained Access

Auditing Admin Access Using CloudTrail

Database Model v3

Database Model v3 Review

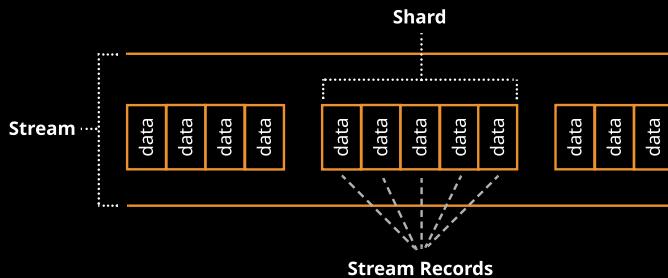
Integration

Amazon EMR

Amazon Elasticsearch Service

What Are DynamoDB Streams?

- Time-ordered sequence of item-level changes in a table
- Partitioned change log (shards) allows scaling
- Information stored for 24 hours
- Provides a stream of inserts, updates, and deletes
- Stream records appear in the same sequence as the item modifications
- Guaranteed to be delivered **only once**
- Use Kinesis Client Library (KCL), Lambda, or the API to query changes
 - **KEYS_ONLY** — Only the **key attributes** of the modified item
 - **NEW_IMAGE** — The entire item, as it appears **after** it was modified
 - **OLD_IMAGE** — The entire item, as it appeared **before** it was modified
 - **NEW_AND_OLD_IMAGES** — Both the **new and the old images** of the item
- Each **GetRecords** API call consumes 1 **streams read request unit** (not RCU) and returns up to 1 MB of data
- The first 2.5 M streams read request units are free; \$0.02 per 100k reads thereafter



Use Cases

- Cross-region replication (global tables)
- Establish relationships across tables
- Messaging/notifications
- Aggregation/filtering
- Analytical reporting
- Archiving and auditing
- Search

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing
High-Performance
Architectures

Static Data Dumps
(streams to static file)

DynamoDB Accelerator
(DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and
Fine-Grained Access

Auditing Admin Access
Using CloudTrail

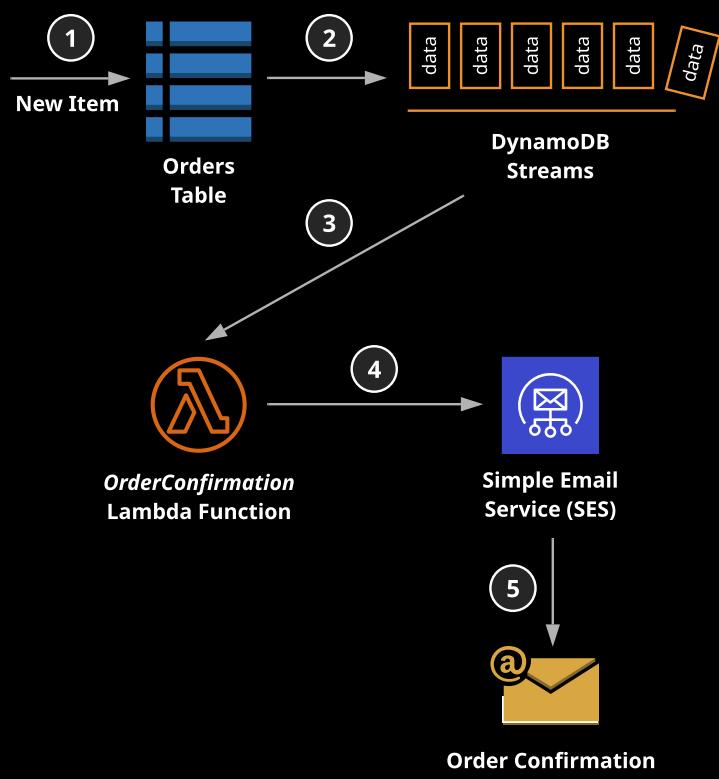
Database Model v3

Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch
Service



Advanced DynamoDB

Advanced Index Usage

Course Navigation

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing
High-Performance
Architectures

Static Data Dumps
(streams to static file)

DynamoDB Accelerator
(DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and
Fine-Grained Access

Auditing Admin Access
Using CloudTrail

Database Model v3

Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch
Service

Select the items below regarding
advanced index usage:

Selective Write Sharding

Aggregation with Streams

Previous Section

Next

Back to Main



Linux Academy

Advanced DynamoDB

Selective Write Sharding

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing
High-Performance
ArchitecturesStatic Data Dumps
(streams to static file)DynamoDB Accelerator
(DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and
Fine-Grained AccessAuditing Admin Access
Using CloudTrail

Database Model v3

Database Model v3 Review

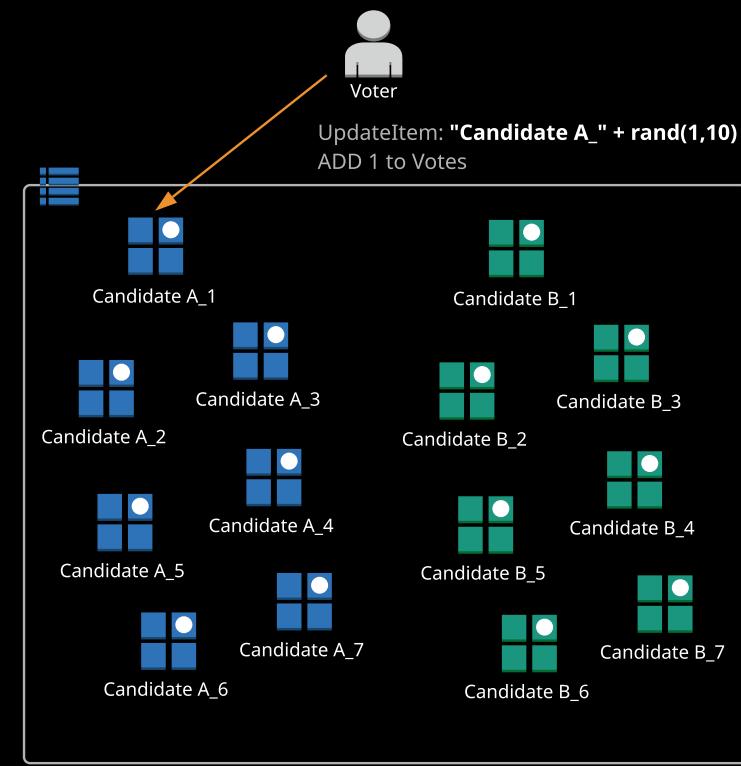
Integration

Amazon EMR

Amazon Elasticsearch
Service

Best Practice: For write-heavy items, avoid hot keys by applying a random suffix to partition keys.

Example: Voting Application

[Next](#)[Back to Main](#)

Linux Academy

Advanced DynamoDB

Selective Write Sharding

Course Navigation

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps
(streams to static file)

DynamoDB Accelerator
(DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and
Fine-Grained Access

Auditing Admin Access
Using CloudTrail

Database Model v3

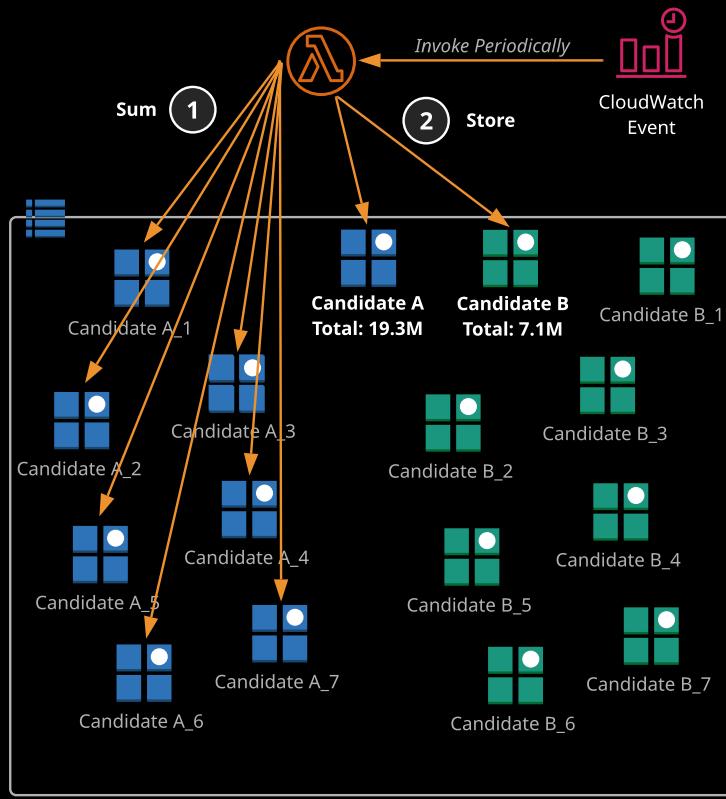
Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch
Service

Aggregation: Use a scheduled Lambda function to query the GSI for the sharded totals, sum them, and write an item back to the table with the aggregate total of votes for the given candidate.



[Back to Main](#)



Linux Academy

Advanced DynamoDB

Aggregation with Streams

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing
High-Performance
Architectures

Static Data Dumps
(streams to static file)

DynamoDB Accelerator
(DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and
Fine-Grained Access

Auditing Admin Access
Using CloudTrail

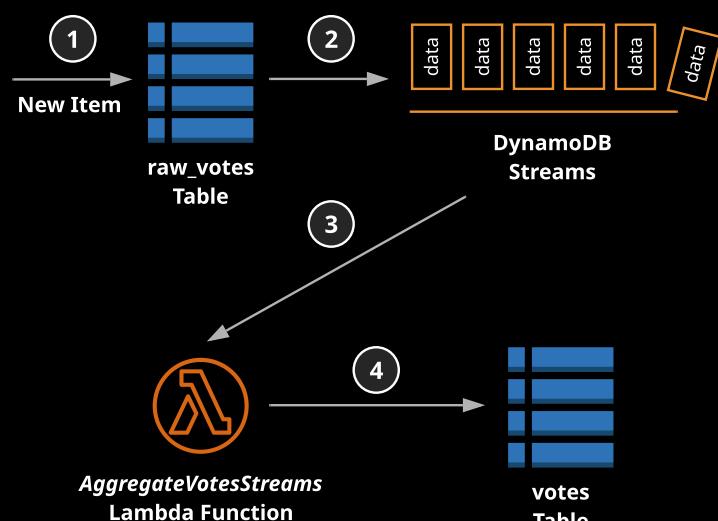
Database Model v3

Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch
Service



The Lambda function will increment the vote count for each candidate's **raw_vote** record in the stream, incrementing the sum to the **votes** table.

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps
(streams to static file)

DynamoDB Accelerator
(DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and
Fine-Grained Access

Auditing Admin Access
Using CloudTrail

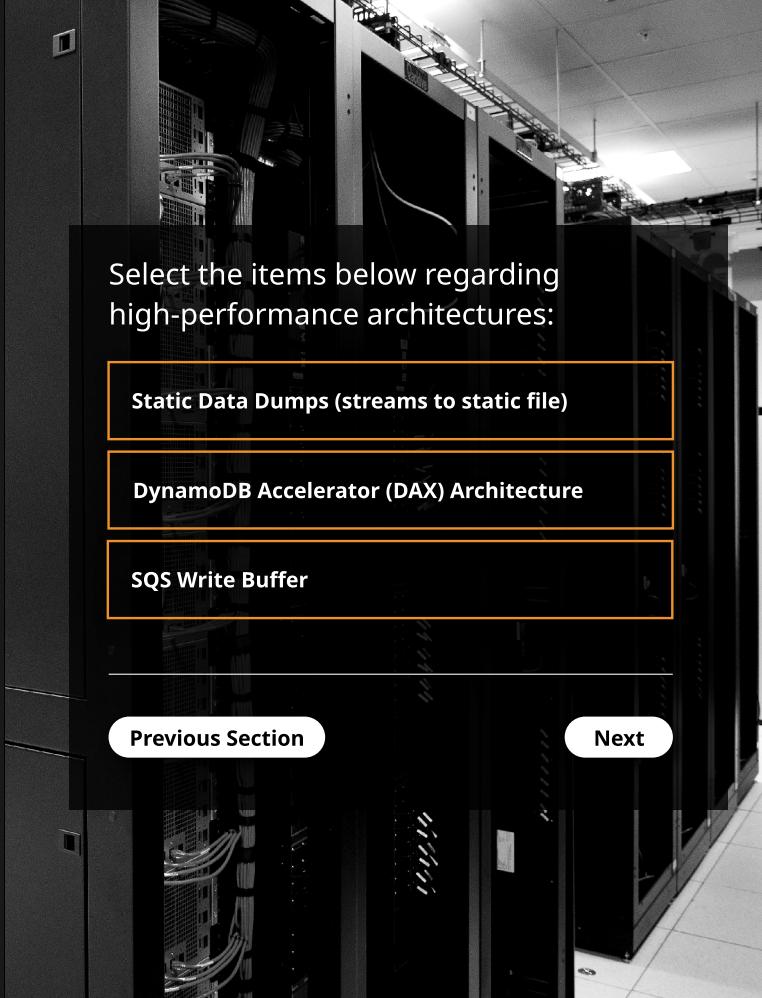
Database Model v3

Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch
Service



Select the items below regarding high-performance architectures:

Static Data Dumps (streams to static file)

DynamoDB Accelerator (DAX) Architecture

SQS Write Buffer

Previous Section

Next

Advanced DynamoDB

Static Data Dumps (streams to static file)

Advanced DynamoDB Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams to static file)

DynamoDB Accelerator (DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and
Fine-Grained Access

Auditing Admin Access
Using CloudTrail

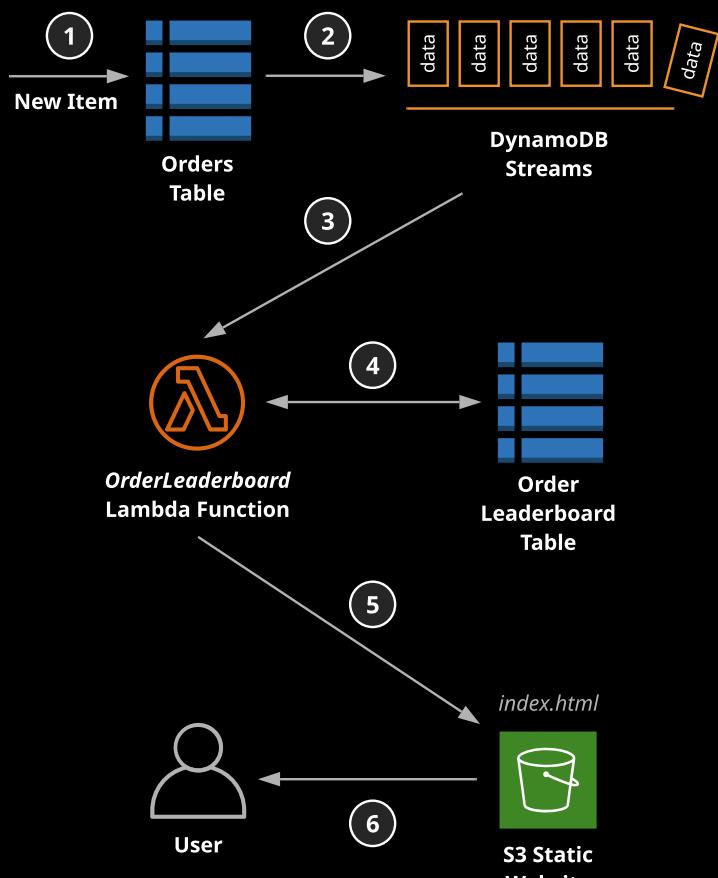
Database Model v3

Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch
Service



Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams to static file)

DynamoDB Accelerator
(DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and
Fine-Grained Access

Auditing Admin Access
Using CloudTrail

Database Model v3

Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch
Service

What Is DAX?

- Fully managed, highly available, in-memory cache for DynamoDB that delivers up to a 10x performance improvement
- Reduces request times from milliseconds to microseconds, even under load
- No need for developers to manage caching logic
- Compatible with DynamoDB API calls

Traditional Cache



Application



Caching with DAX



Application



Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams to static file)

DynamoDB Accelerator (DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and
Fine-Grained AccessAuditing Admin Access
Using CloudTrail

Database Model v3

Database Model v3 Review

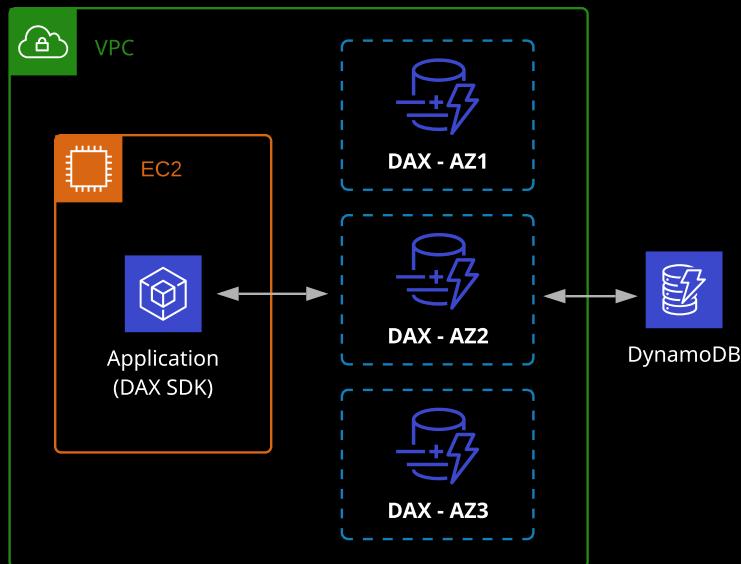
Integration

Amazon EMR

Amazon Elasticsearch
Service

High Availability

DAX is designed to run within a VPC. In the event of a primary node failure, DAX fails over to a read replica and promotes that node as the new primary.



API Compatibility

DAX is API compatible with DynamoDB.

- **Read APIs:** GetItem, BatchGetItem, Query, Scan
- **Modify APIs:** PutItem, UpdateItem, DeleteItem, BatchWriteItem
- **Control plane APIs:** Not supported (CreateTable, DeleteTable, etc.)

[Back](#)[Back to Main](#)

Linux Academy

Advanced DynamoDB

SQS Write Buffer

Course Navigation

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams to static file)

DynamoDB Accelerator (DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and Fine-Grained Access

Auditing Admin Access Using CloudTrail

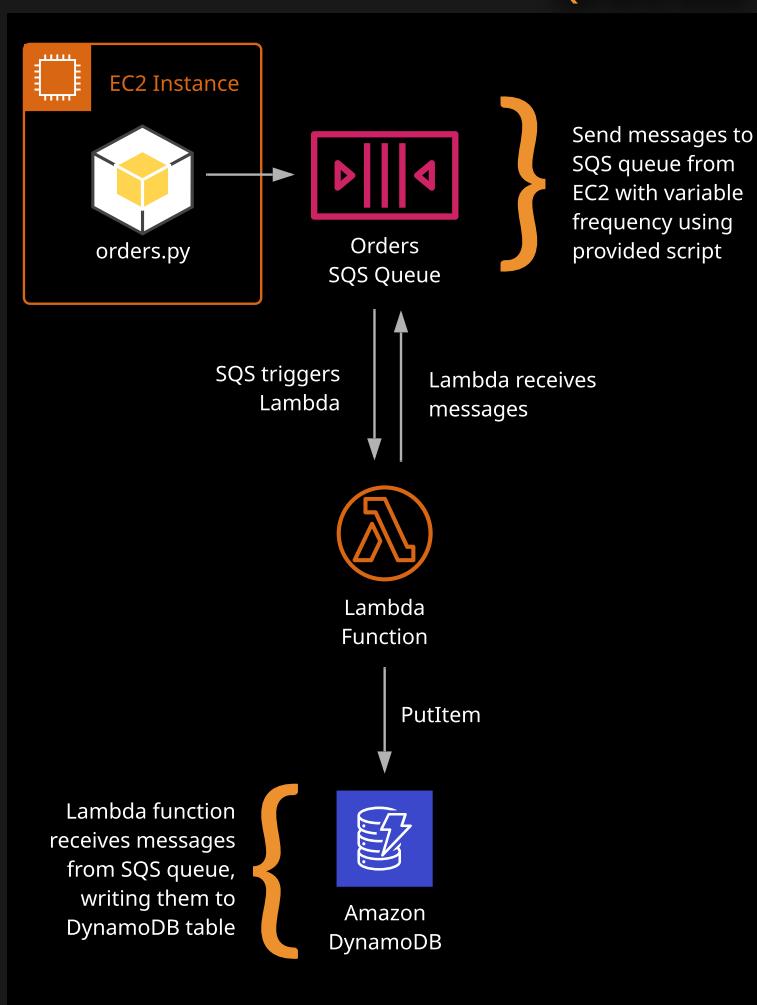
Database Model v3

Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch Service



[Back to Main](#)



Linux Academy

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams to static file)

DynamoDB Accelerator (DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and Fine-Grained Access

Auditing Admin Access Using CloudTrail

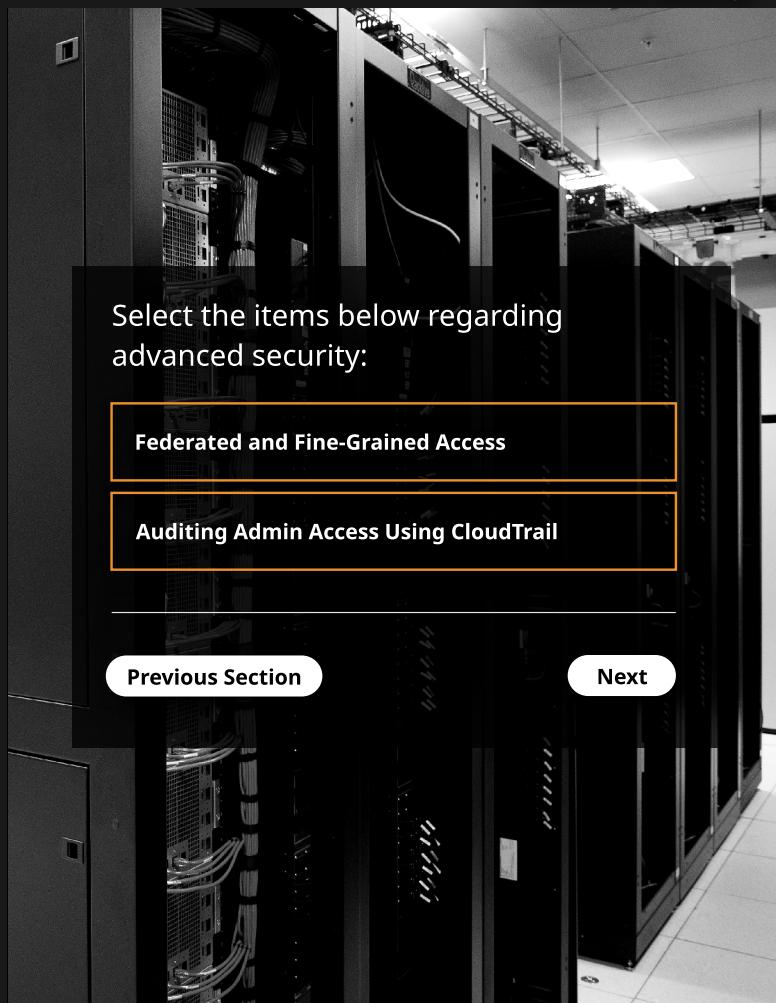
Database Model v3

Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch Service



Select the items below regarding advanced security:

Federated and Fine-Grained Access

Auditing Admin Access Using CloudTrail

[Previous Section](#)

[Next](#)

Advanced DynamoDB

Federated and Fine-Grained Access

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams to static file)

DynamoDB Accelerator (DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and Fine-Grained Access

Auditing Admin Access
Using CloudTrail

Database Model v3

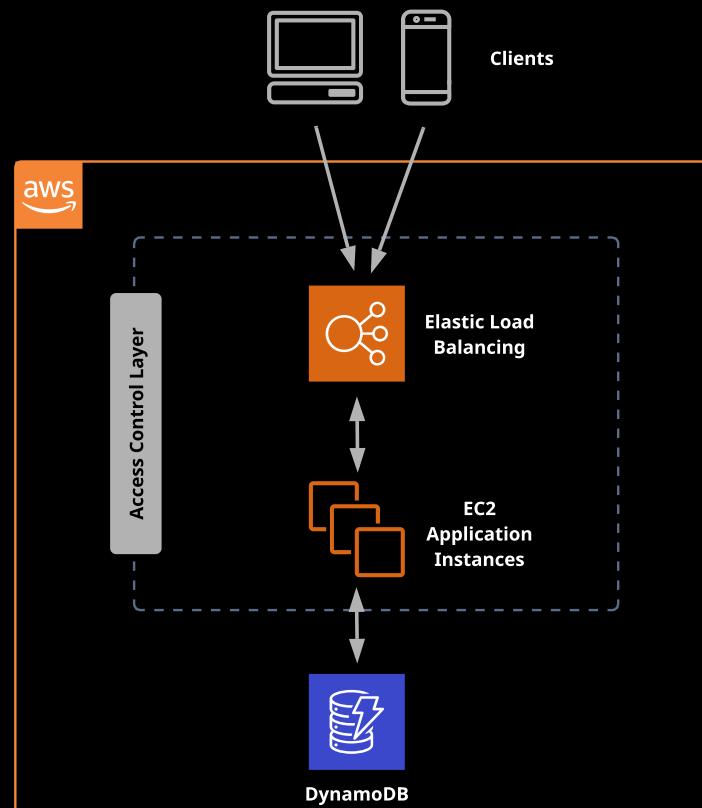
Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch Service

Without Fine-Grained Access Control



Next

Back to Main



Linux Academy

Advanced DynamoDB

Federated and Fine-Grained Access

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams to static file)

DynamoDB Accelerator (DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and Fine-Grained Access

Auditing Admin Access
Using CloudTrail

Database Model v3

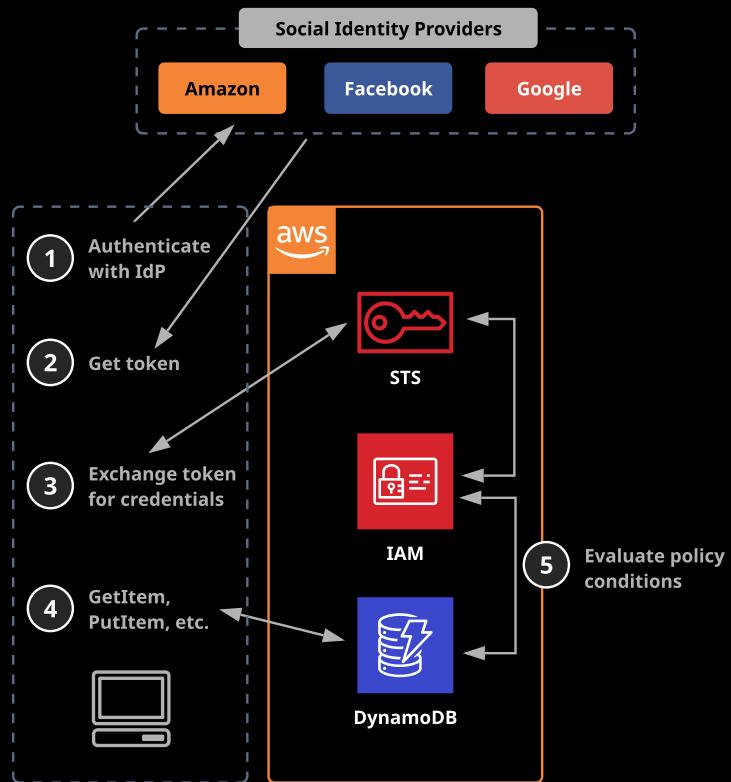
Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch Service

With Fine-Grained Access Control

[Back](#)[Next](#)[Back to Main](#)

Linux Academy

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams
to static file)DynamoDB Accelerator
(DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and Fine-Grained Access

Auditing Admin Access
Using CloudTrail

Database Model v3

Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch
Service

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Sid": "FullAccessToUserItems",  
            "Effect": "Allow",  
            "Action": [  
                "dynamodb:GetItem",  
                "dynamodb:BatchGetItem",  
                "dynamodb:Query",  
                "dynamodb:PutItem",  
                "dynamodb:UpdateItem",  
                "dynamodb:DeleteItem",  
                "dynamodb:BatchWriteItem"  
            ],  
            "Resource": [  
                "arn:aws:dynamodb:us-east-1:123456789012:table/MyTable"  
            ],  
            "Condition": {  
                "ForAllValues:StringEquals": {  
                    "dynamodb:LeadingKeys": [  
                        "${www.amazon.com:user_id}"  
                    ]  
                }  
            }  
        }  
    ]  
}
```

Advanced DynamoDB

Section 6

Streams and Triggers

[Streams](#)[Triggers \(Lambda Invoke\)](#)

Advanced Index Usage

[Selective Write Sharding](#)[Aggregation with Streams](#)

Implementing High-Performance Architectures

[Static Data Dumps \(streams to static file\)](#)[DynamoDB Accelerator \(DAX\) Architecture](#)[SQS Write Buffer](#)

Advanced Security

Federated and Fine-Grained Access

[Auditing Admin Access](#)[Using CloudTrail](#)

Database Model v3

[Database Model v3 Review](#)

Integration

[Amazon EMR](#)[Amazon Elasticsearch Service](#)

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Sid": "AllowAccessToOnlyItemsMatchingUserID",  
            "Effect": "Allow",  
            "Action": [  
                "dynamodb:GetItem",  
                "dynamodb:BatchGetItem",  
                "dynamodb:Query",  
                "dynamodb:PutItem",  
                "dynamodb:UpdateItem",  
                "dynamodb:DeleteItem",  
                "dynamodb:BatchWriteItem"  
            ],  
            "Resource": [  
                "arn:aws:dynamodb:us-east-1:123456789012:table/MyTable"  
            ],  
            "Condition": {  
                "ForAllValues:StringEquals": {  
                    "dynamodb:LeadingKeys": [  
                        "${www.amazon.com:user_id}"  
                    ],  
                    "dynamodb:Attributes": [  
                        "userid",  
                        "artist_name",  
                        "album_title",  
                        "year",  
                        "tracks"  
                    ]  
                },  
                "StringEqualsIfExists": {  
                    "dynamodb:Select": "SPECIFIC_ATTRIBUTES"  
                }  
            }  
        }  
    ]  
}
```



Advanced DynamoDB

Section 6

Streams and Triggers

- Streams
- Triggers (Lambda Invoke)

Advanced Index Usage

- Selective Write Sharding
- Aggregation with Streams

Implementing High-Performance Architectures

- Static Data Dumps (streams to static file)
- DynamoDB Accelerator (DAX) Architecture
- SQS Write Buffer

Advanced Security

Federated and Fine-Grained Access

- Auditing Admin Access
- Using CloudTrail

Database Model v3

- Database Model v3 Review

Integration

- Amazon EMR
- Amazon Elasticsearch Service

Advanced DynamoDB

Federated and Fine-Grained Access

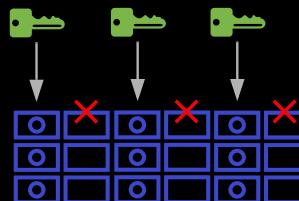
IAM Policy Conditions

Control access to individual data items and attributes using IAM policies.



Grant permissions to a table, but restrict items based on PK values.

[Sample Policy](#)



Hide information so that only a subset of attributes is visible to the user.

[Sample Policy](#)

[Back](#)

[Back to Main](#)



Linux Academy

Advanced DynamoDB

Section 6

Streams and Triggers

- Streams
- Triggers (Lambda Invoke)

Advanced Index Usage

- Selective Write Sharding
- Aggregation with Streams

Implementing High-Performance Architectures

- Static Data Dumps (streams to static file)
- DynamoDB Accelerator (DAX) Architecture
- SQS Write Buffer

Advanced Security

- Federated and Fine-Grained Access
- Auditing Admin Access Using CloudTrail

Database Model v3

- Database Model v3 Review

Integration

- Amazon EMR
- Amazon Elasticsearch Service

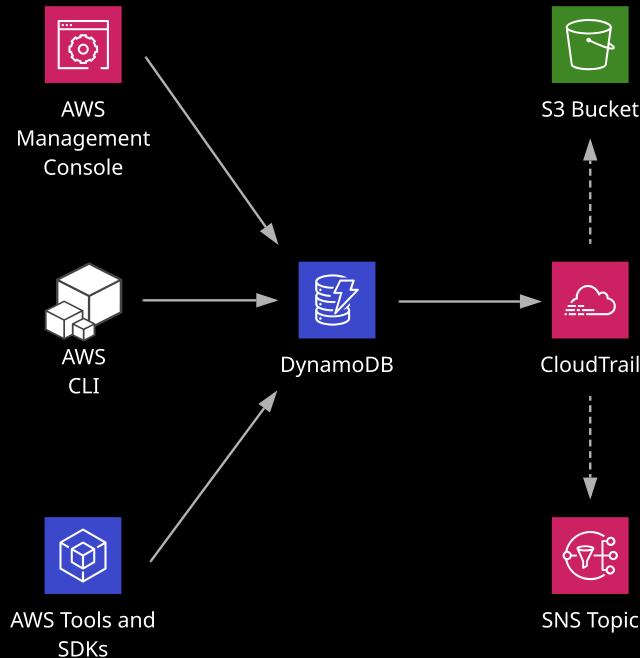
Advanced DynamoDB

Auditing Admin Access Using CloudTrail

CloudTrail provides a record of actions taken by a user, role, or an AWS service in DynamoDB. CloudTrail captures all API calls for DynamoDB as events.

Most DynamoDB, DynamoDB Streams, and DAX API calls are captured by CloudTrail, and optionally stored in an S3 bucket.

SNS can notify you when CloudTrail publishes new log files to your S3 bucket.



Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams
to static file)

DynamoDB Accelerator
(DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and Fine-Grained
Access

Auditing Admin Access
Using CloudTrail

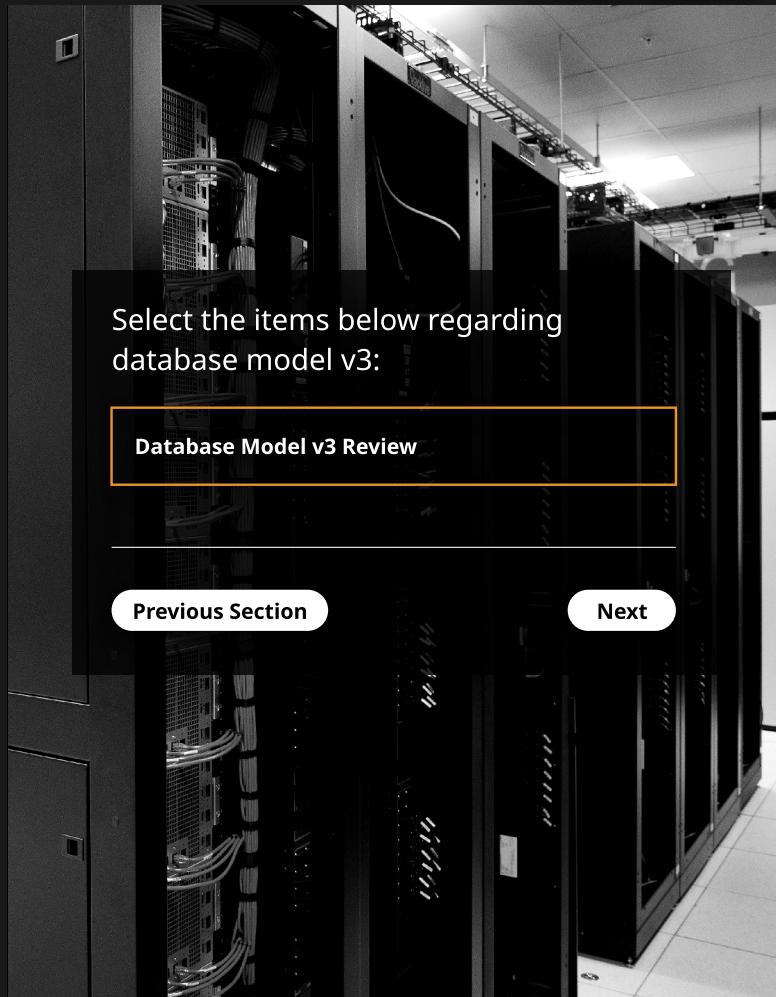
Database Model v3

Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch
Service



Advanced DynamoDB

Database Model v3 Review

Course Navigation

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams to static file)

DynamoDB Accelerator (DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and Fine-Grained Access

Auditing Admin Access Using CloudTrail

Database Model v3

Database Model v3 Review

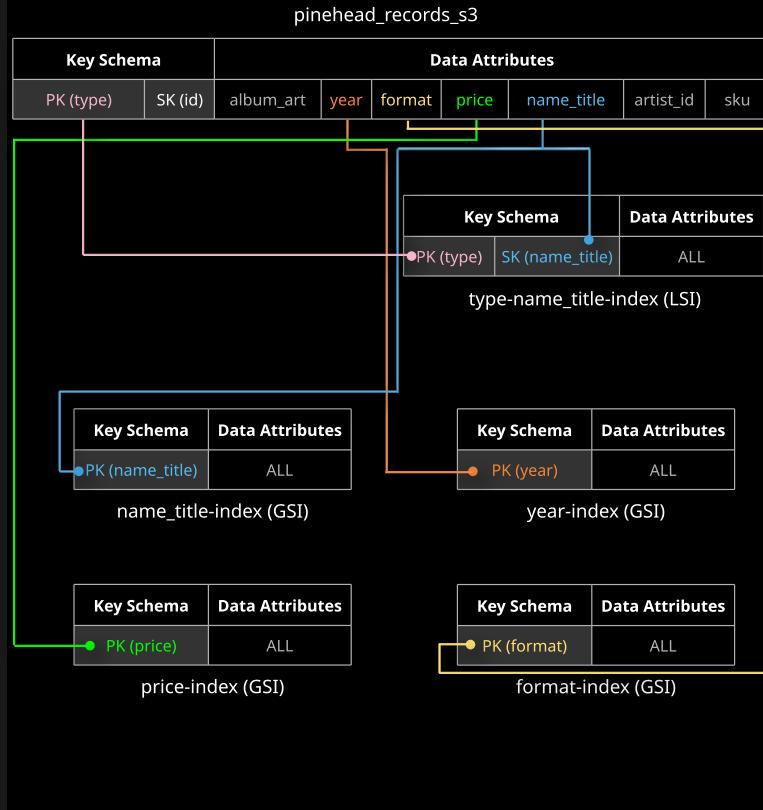
Integration

Amazon EMR

Amazon Elasticsearch Service

This version of the Pinehead Records data model features:

- Entity type partition key schema, where **type** is one of: album, artist, track
- More flexible schema to handle additional entity types



[Back to Main](#)



Linux Academy

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams to static file)

DynamoDB Accelerator (DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and Fine-Grained Access

Auditing Admin Access
Using CloudTrail

Database Model v3

Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch Service

Select the items below regarding integration:

Amazon EMR

Amazon Elasticsearch Service

Previous Section

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams to static file)

DynamoDB Accelerator (DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and Fine-Grained Access

Auditing Admin Access Using CloudTrail

Database Model v3

Database Model v3 Review

Integration

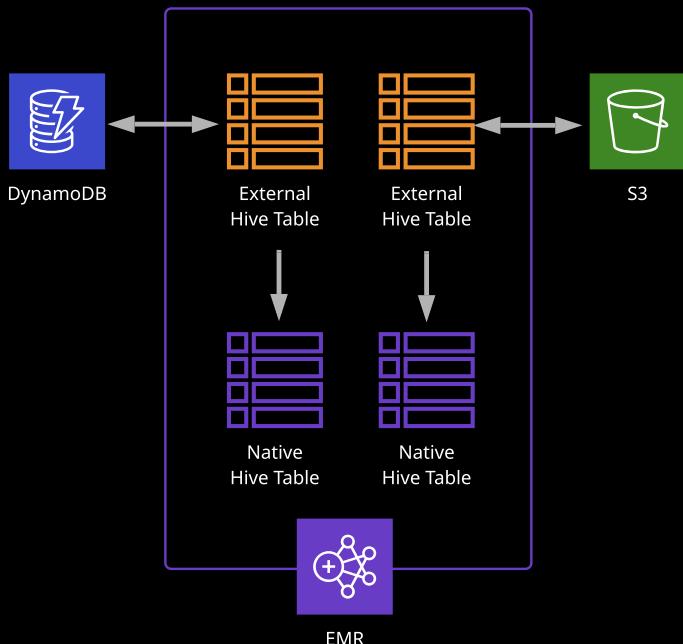
Amazon EMR

Amazon Elasticsearch Service

Using **Amazon EMR** and **Apache Hive**, you can quickly and efficiently process large amounts of data in DynamoDB. Hive allows you to query **MapReduce** clusters using a SQL-like language called **HiveQL**.

EMR comes included with a version of Hive that supports DynamoDB operations such as:

- Loading DynamoDB data into an EMR cluster
- Querying live DynamoDB data using HiveQL
- Exporting data in DynamoDB to S3; importing data in S3 to DynamoDB
- Joining data in DynamoDB with data in S3



Advanced DynamoDB

Amazon Elasticsearch Service

Course Navigation

Advanced DynamoDB

Section 6

Streams and Triggers

Streams

Triggers (Lambda Invoke)

Advanced Index Usage

Selective Write Sharding

Aggregation with Streams

Implementing High-Performance Architectures

Static Data Dumps (streams to static file)

DynamoDB Accelerator (DAX) Architecture

SQS Write Buffer

Advanced Security

Federated and Fine-Grained Access

Auditing Admin Access Using CloudTrail

Database Model v3

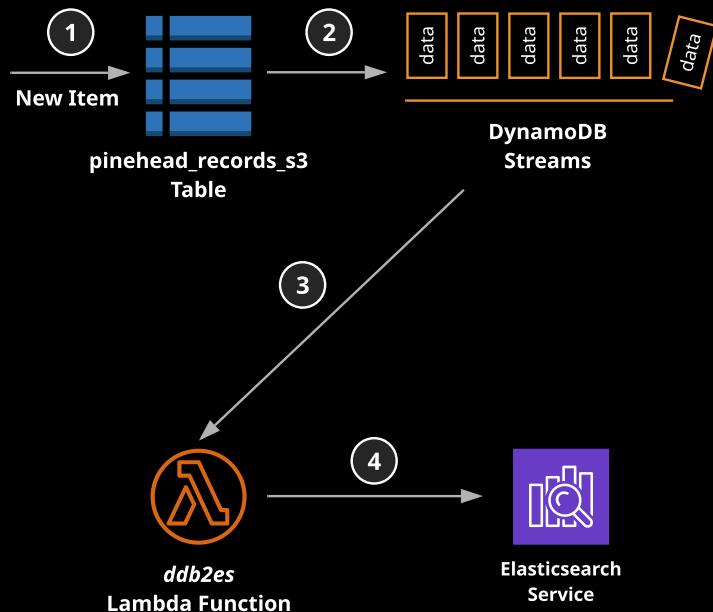
Database Model v3 Review

Integration

Amazon EMR

Amazon Elasticsearch Service

You can use Lambda to send data to your Elasticsearch domain from DynamoDB. New data that arrives in the table triggers an event notification to a Lambda function, which then performs the indexing.



[Back to Main](#)



Linux Academy